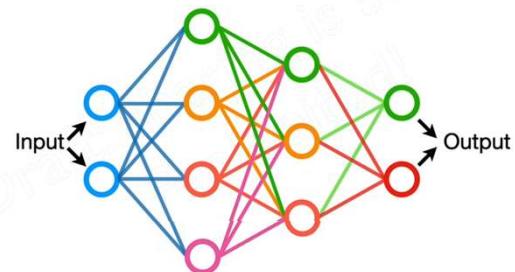
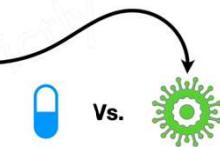


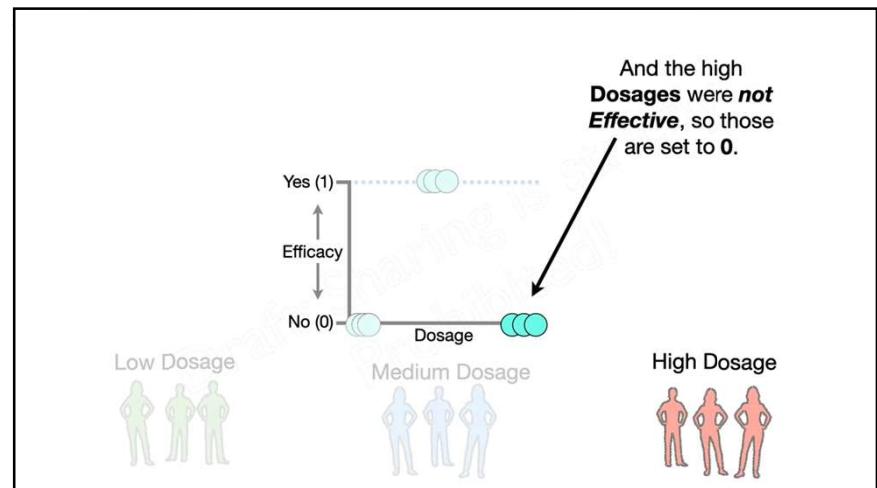
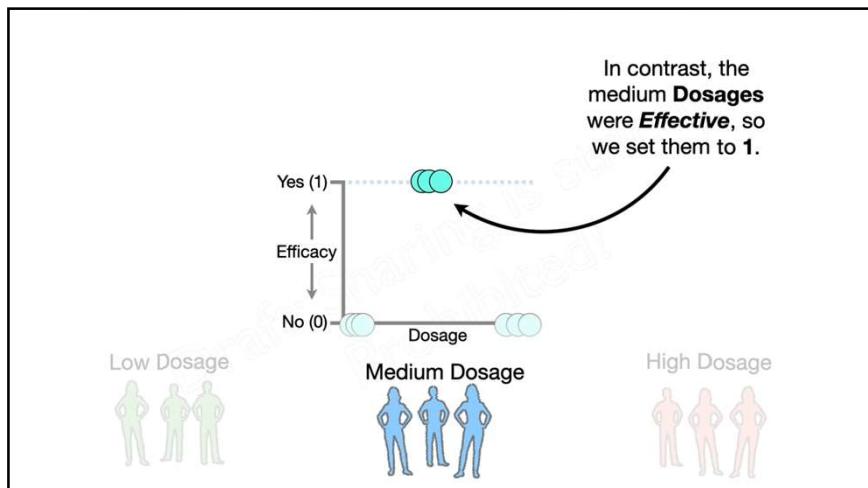
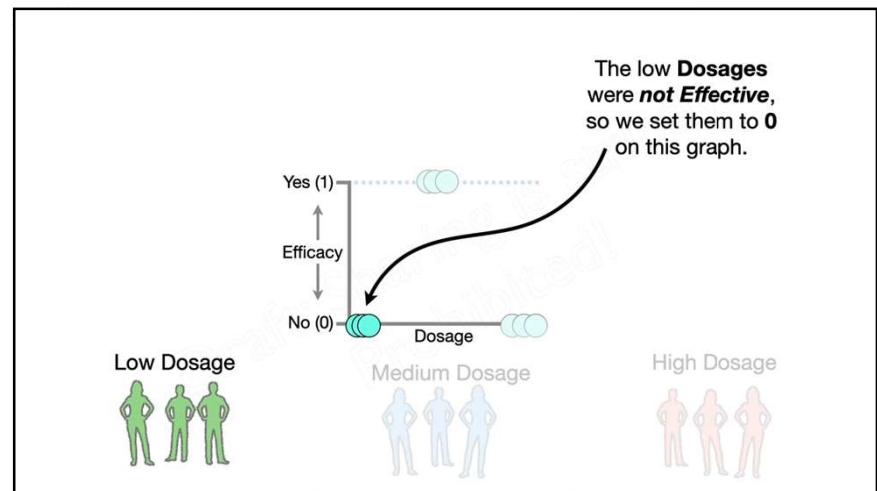
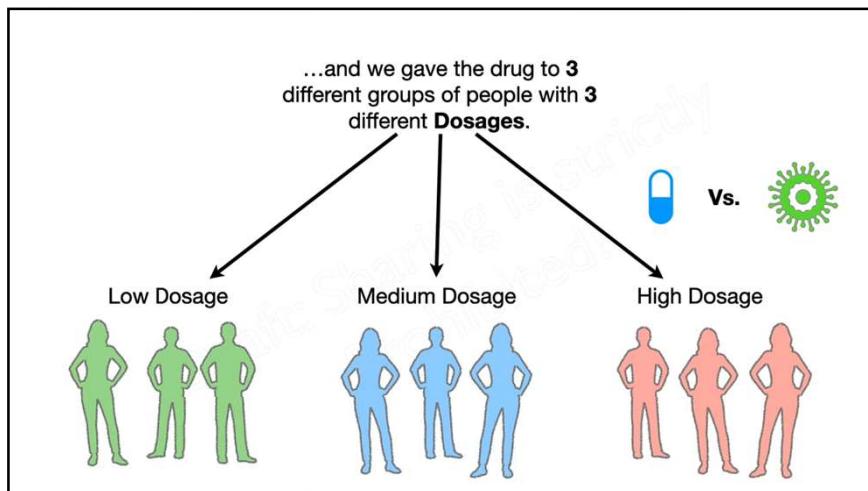
Neural Networks...

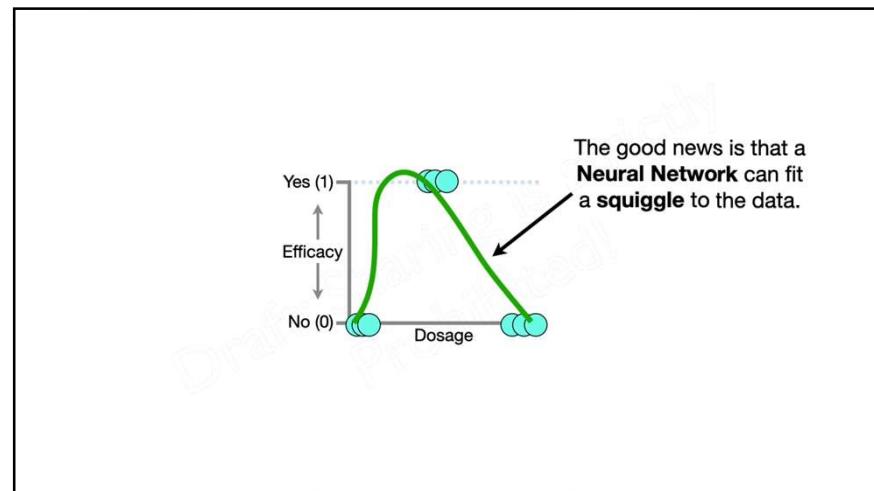
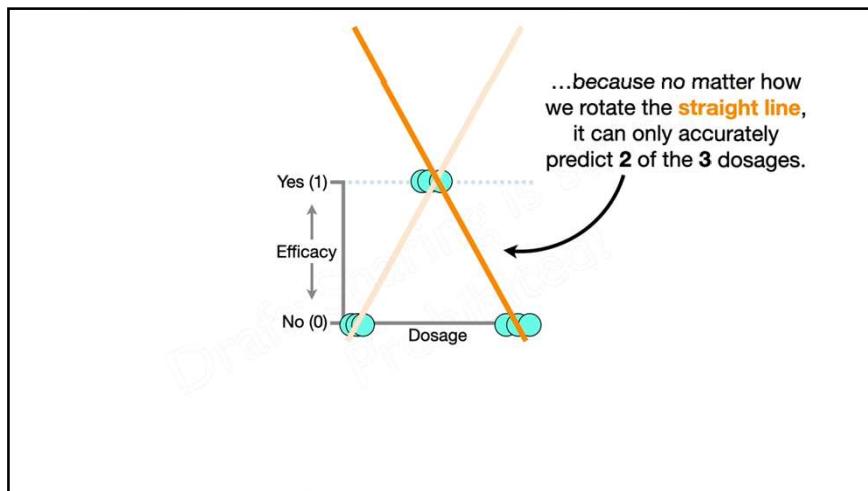
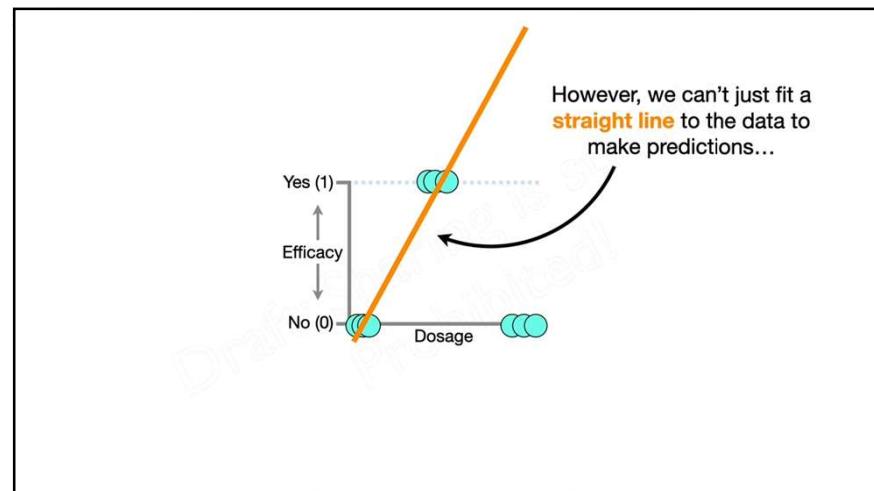
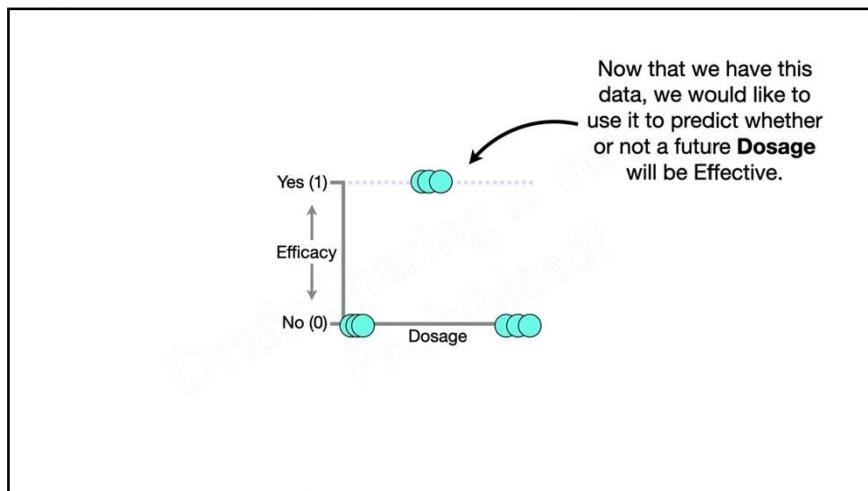
Neural Networks, one of the most popular algorithms in **Machine Learning**, cover a broad range of concepts and techniques.

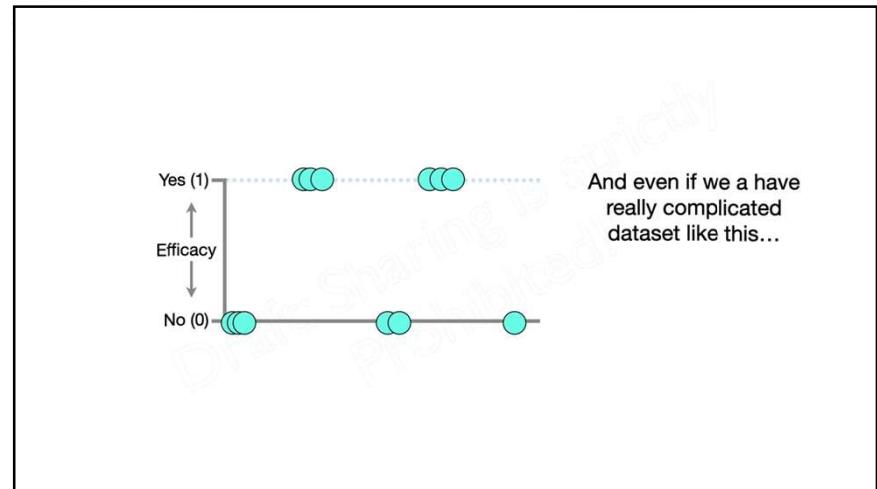
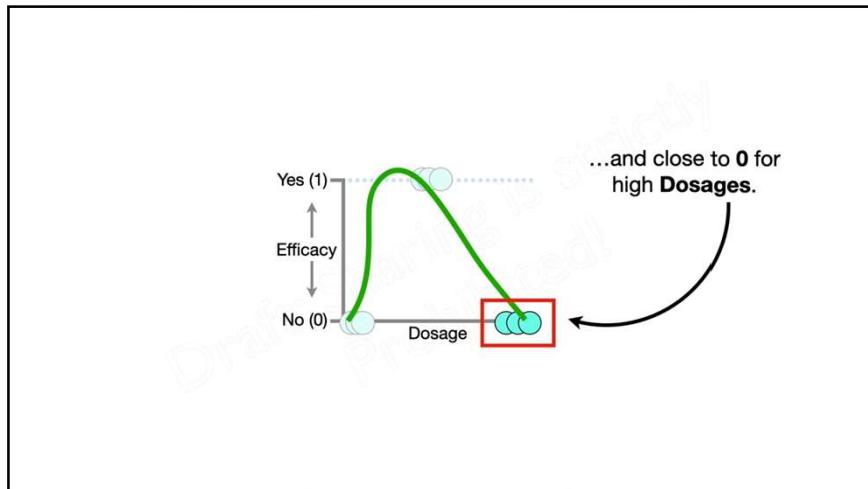
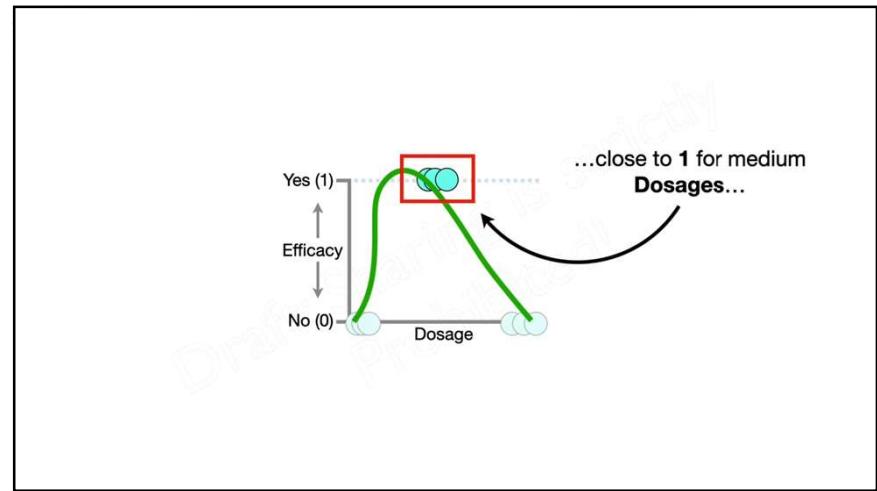
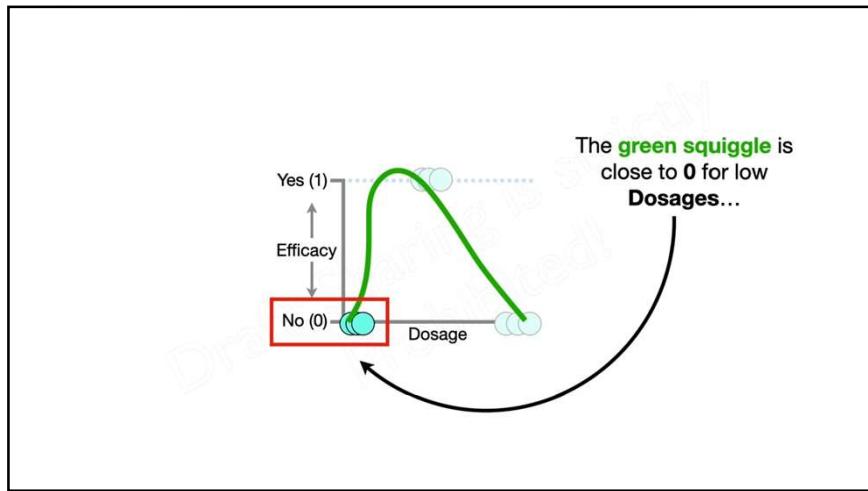


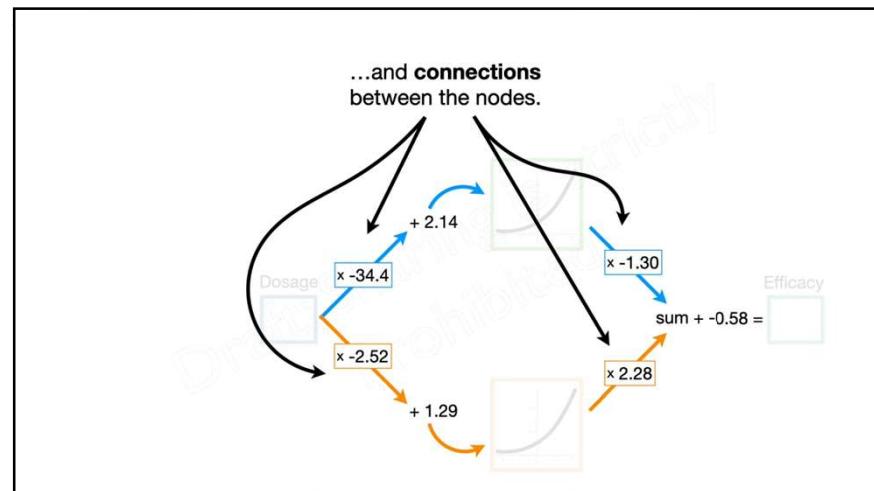
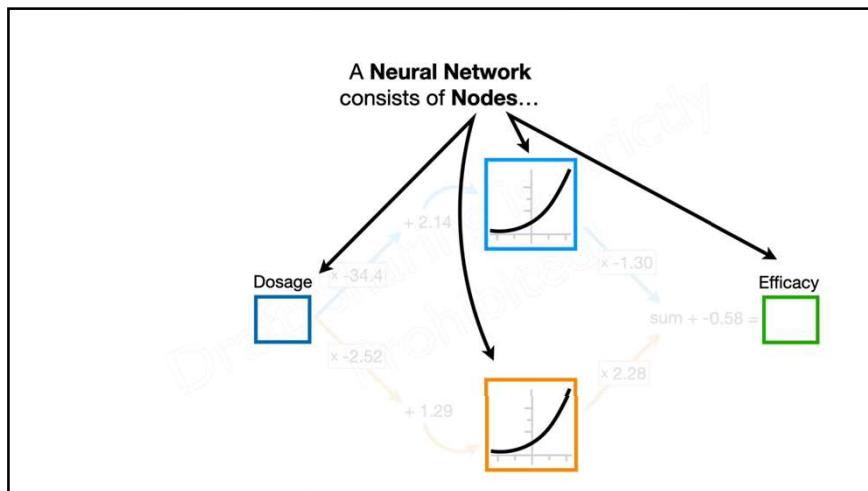
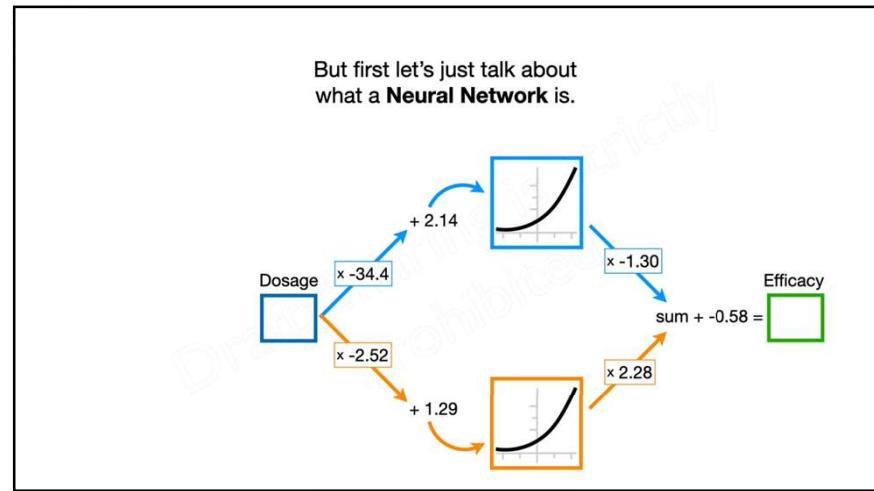
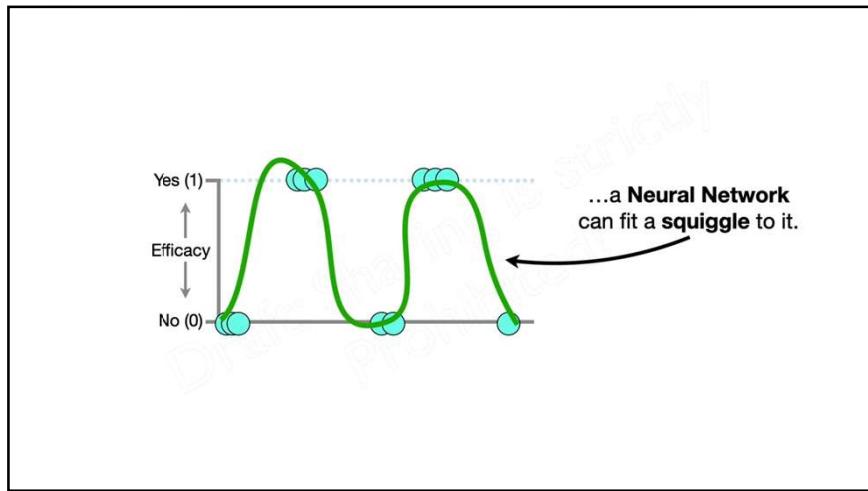
So, with that said, let's imagine we tested a drug that was designed to treat an illness...



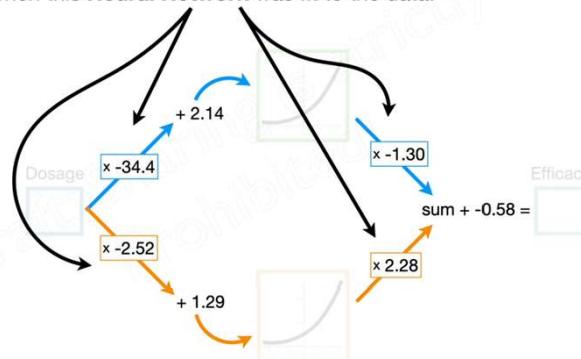




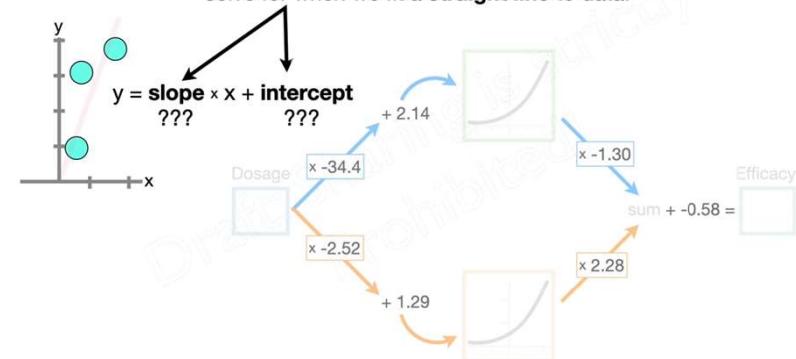




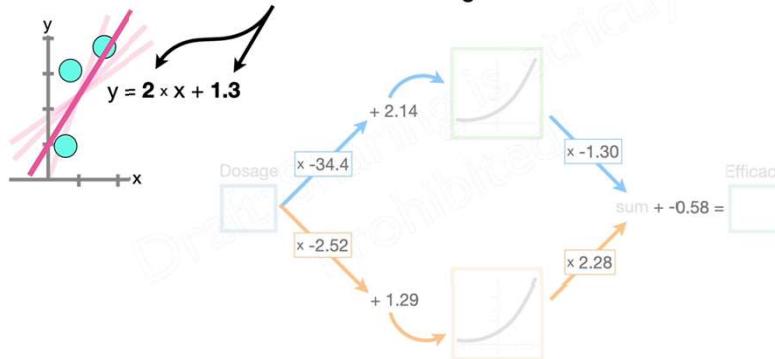
NOTE: The numbers along each connection represent parameter values that were estimated when this **Neural Network** was fit to the data.



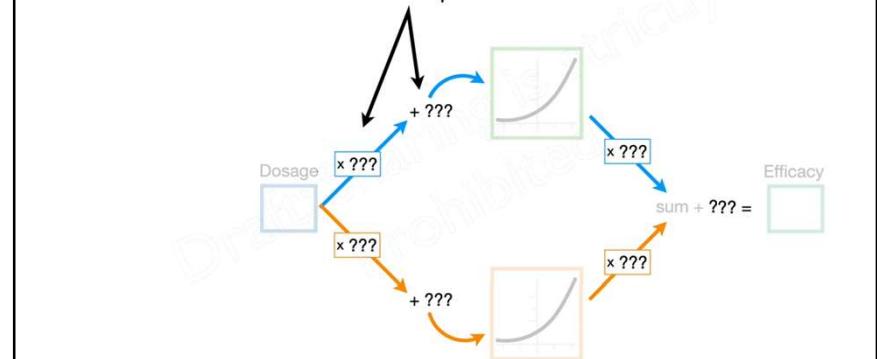
For now, just know that these parameter estimates are analogous to the **slope** and **intercept** values that we solve for when we fit a **straight line** to data.



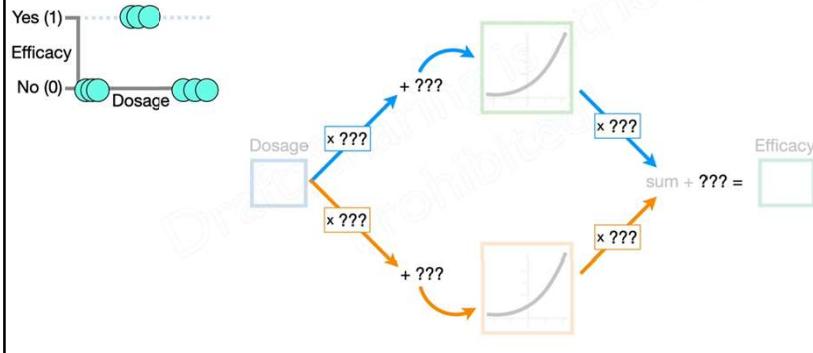
For now, just know that these parameter estimates are analogous to the **slope** and **intercept** values that we solve for when we fit a **straight line** to data.



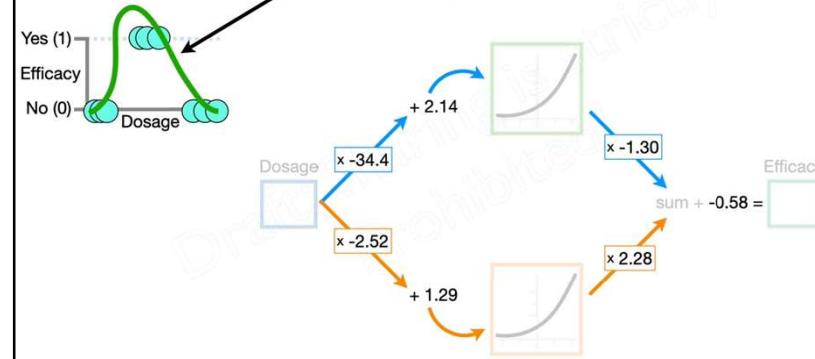
Likewise, a **Neural Network** starts out with unknown parameter values...



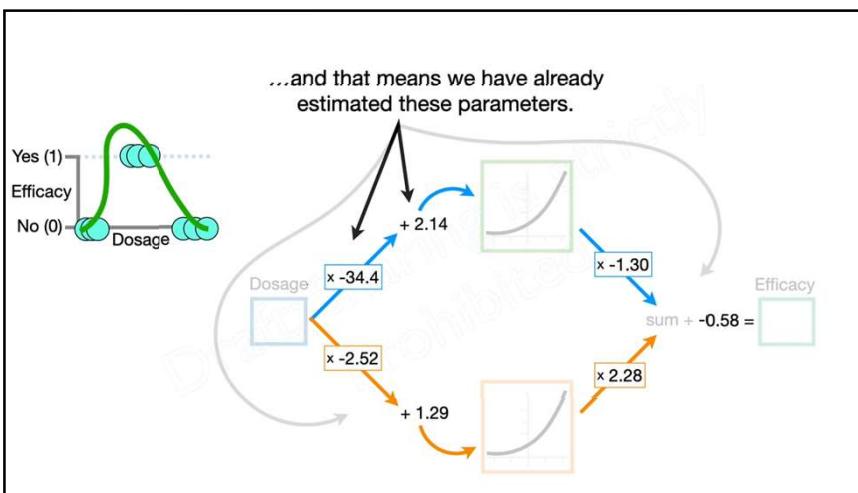
...that are estimated when we fit the **Neural Network** to a dataset using a method called **Backpropagation**.



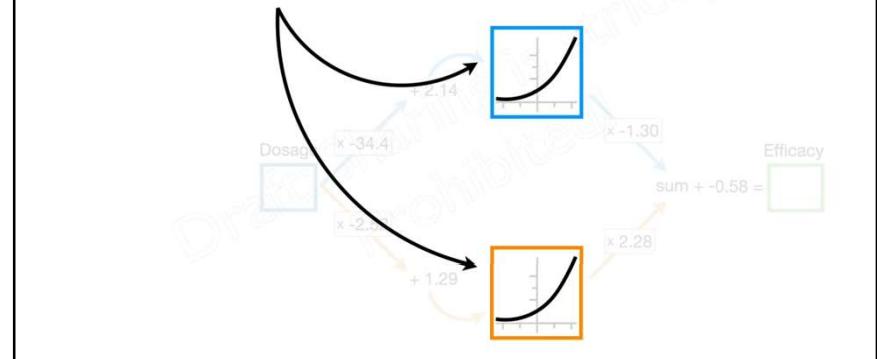
...but for now, just assume that we've already fit this **Neural Network** to this specific dataset...

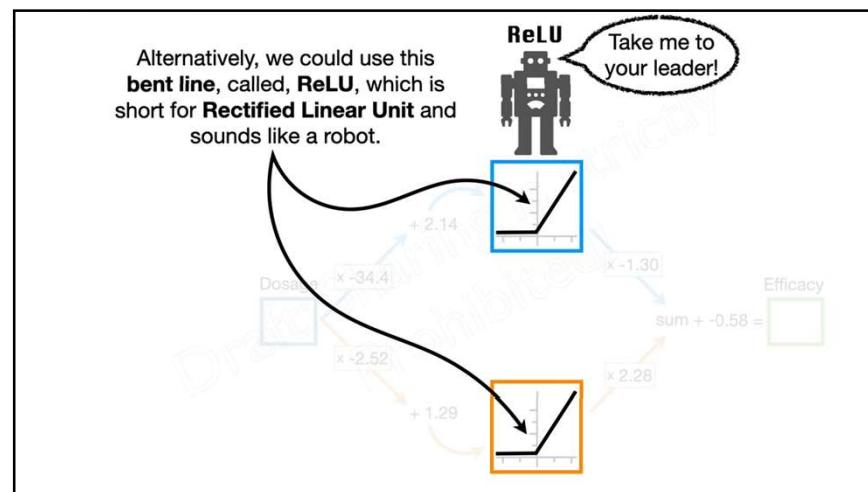
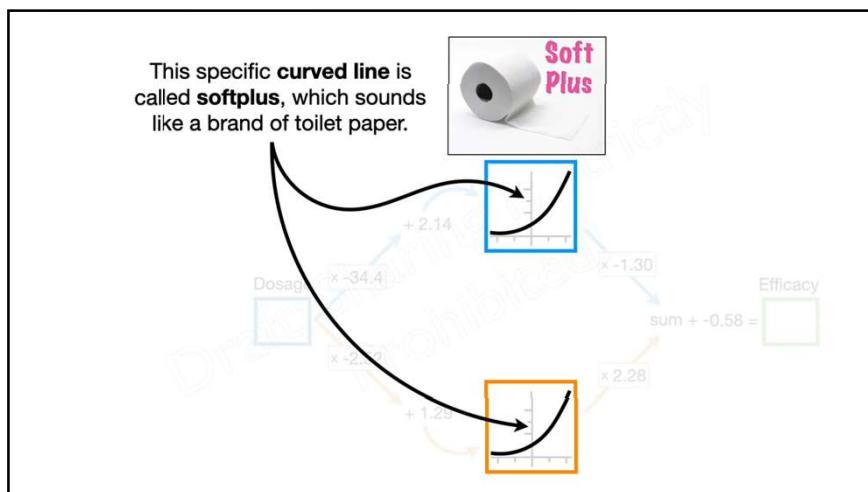
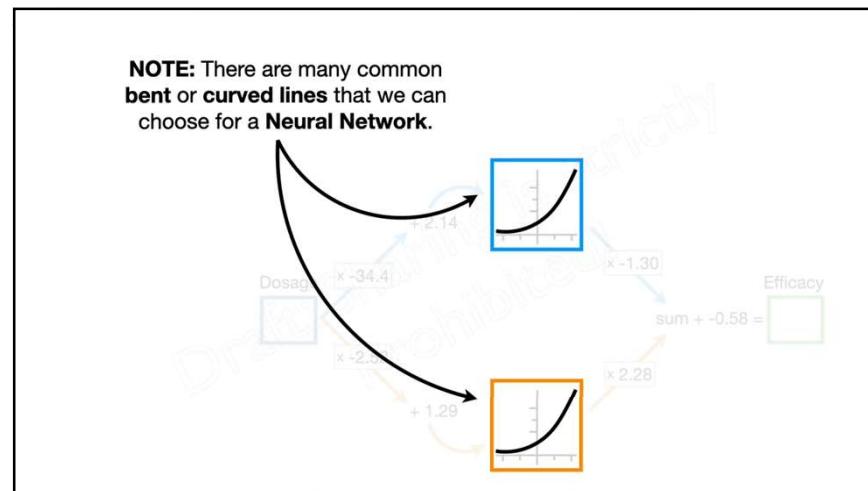
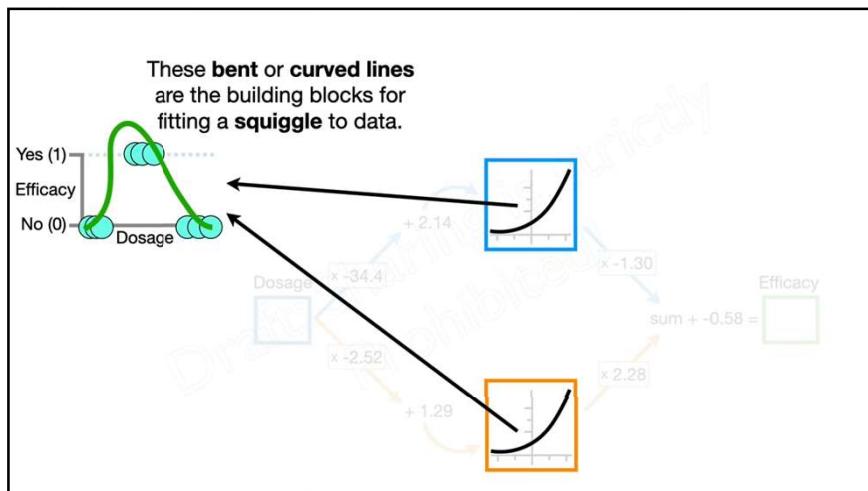


...and that means we have already estimated these parameters.

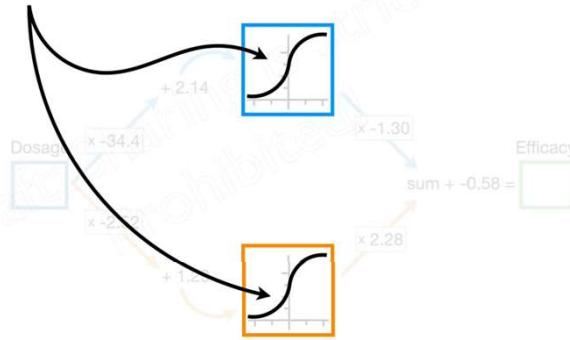


Also, you may have noticed that some of the **Nodes** have **curved lines** inside of them.

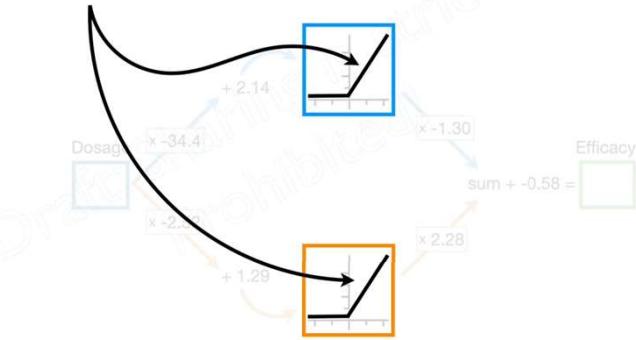




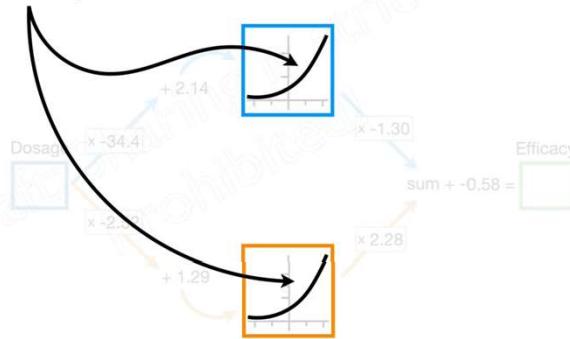
Or we could use a **sigmoid** shape or any other **bent** or **curved line**.



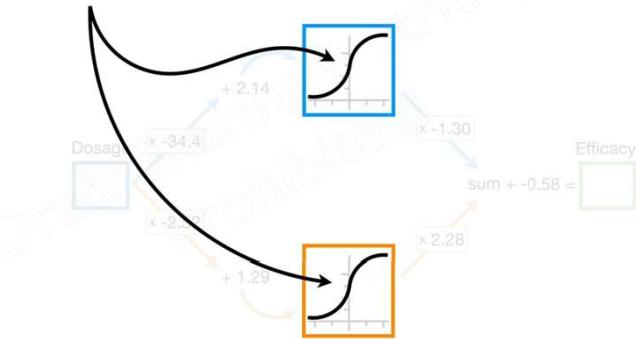
The **curved or bent lines** are called **Activation Functions**.



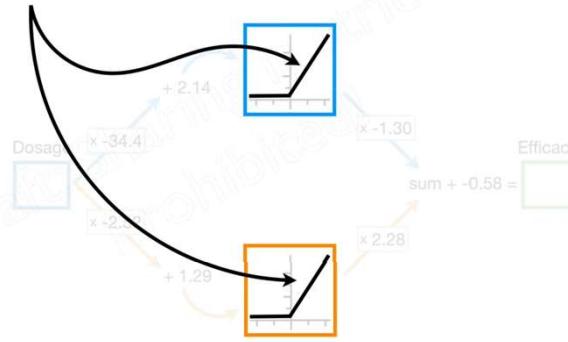
When you build a **Neural Network**, you have to decide which **Activation Function or Functions** you want to use.



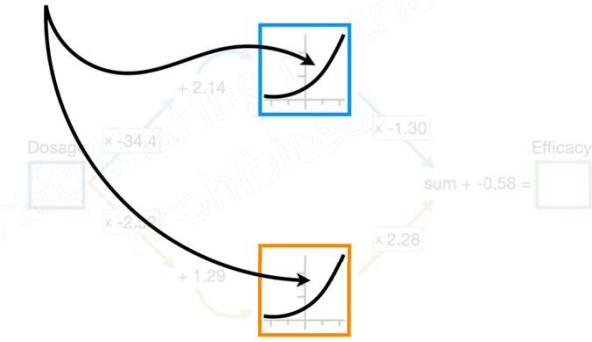
When most people **teach Neural Networks**, they use the **sigmoid Activation Function**.



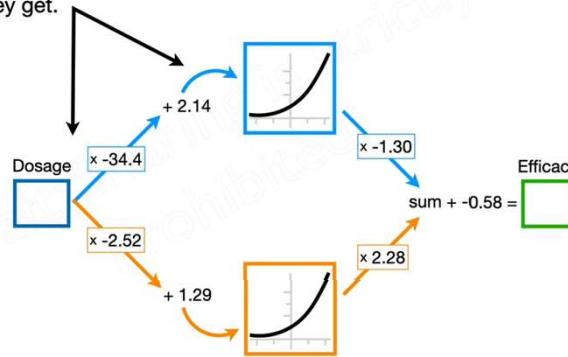
However, in *practice*, it is much more common to use the **ReLU Activation Function**...



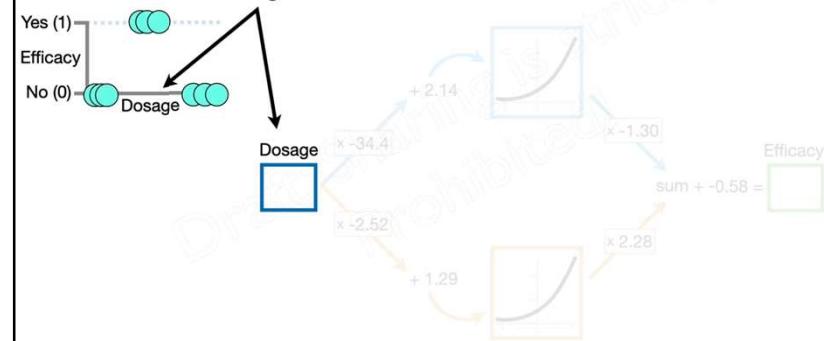
...or the **softplus Activation Function**.

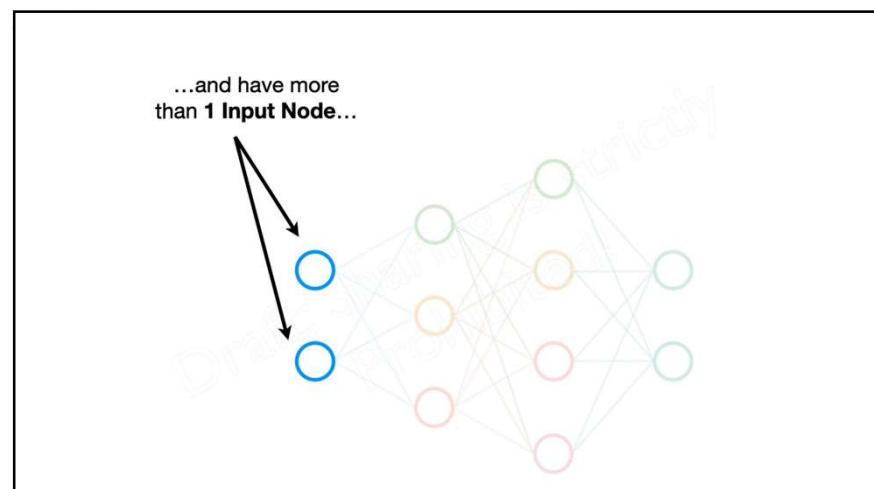
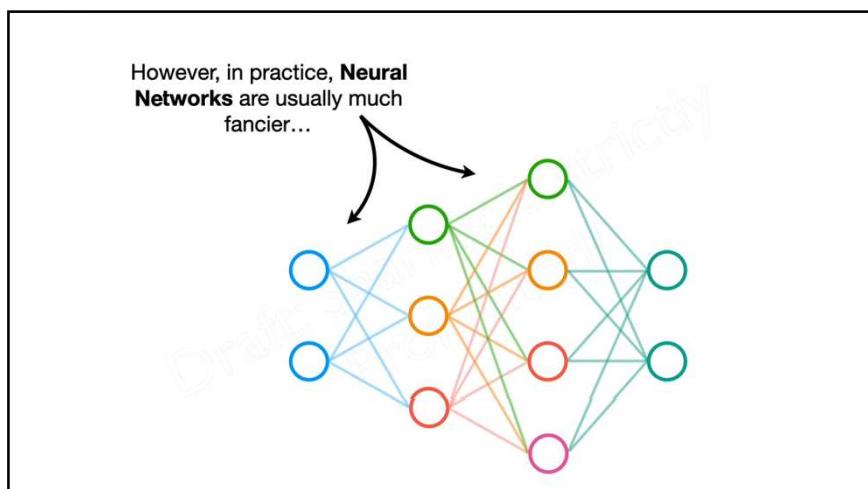
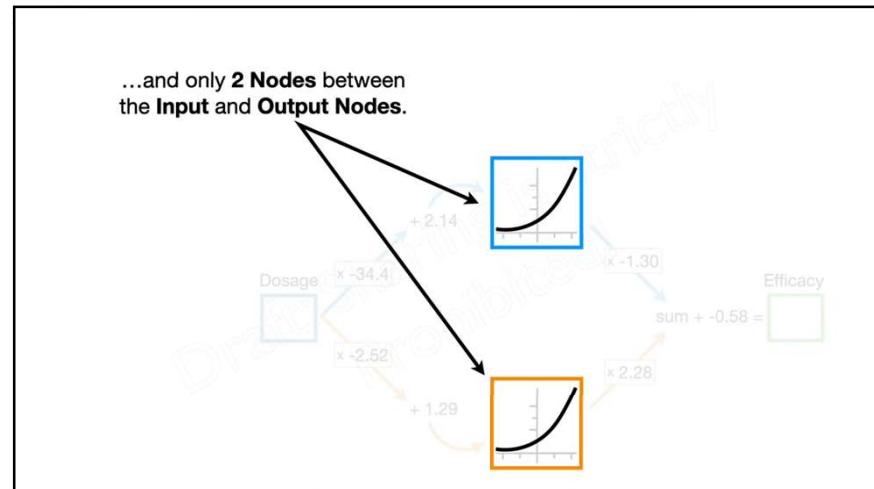
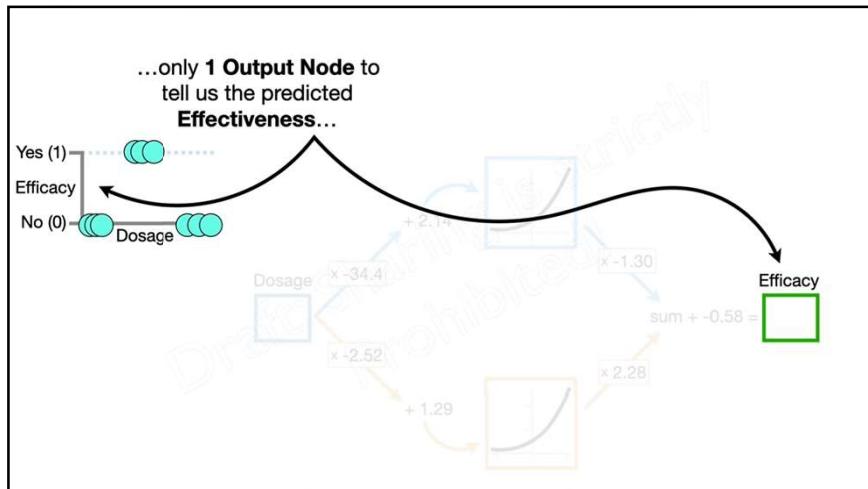


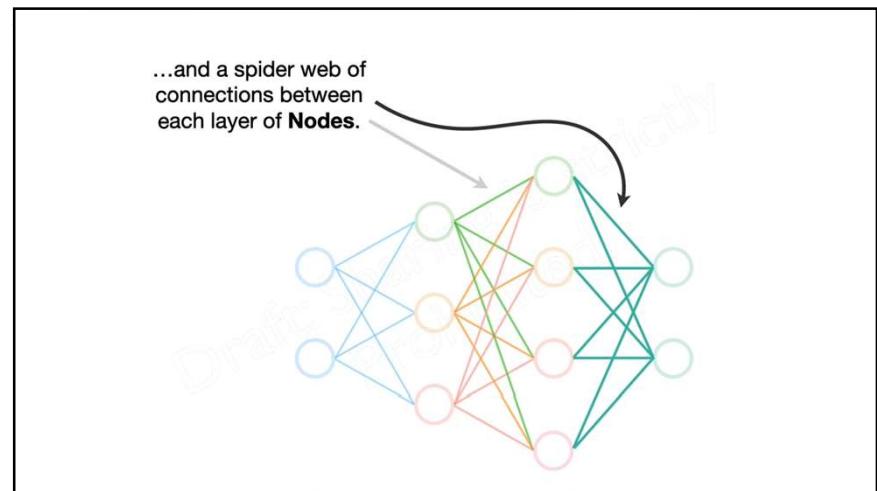
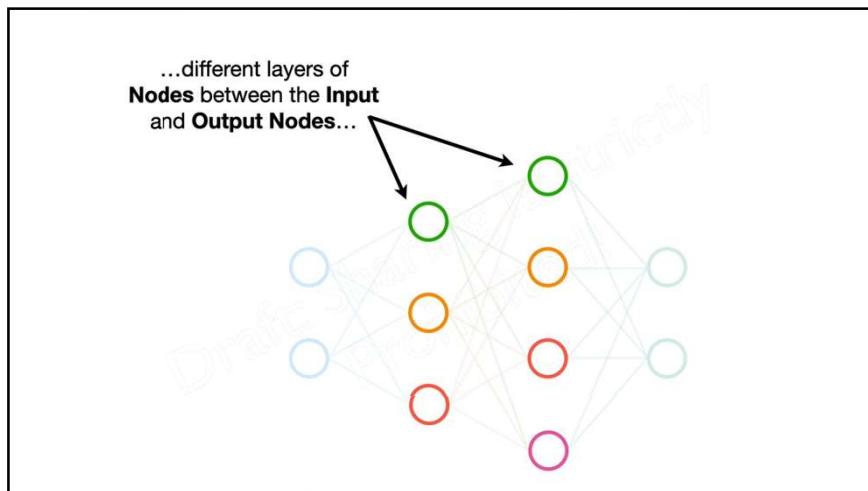
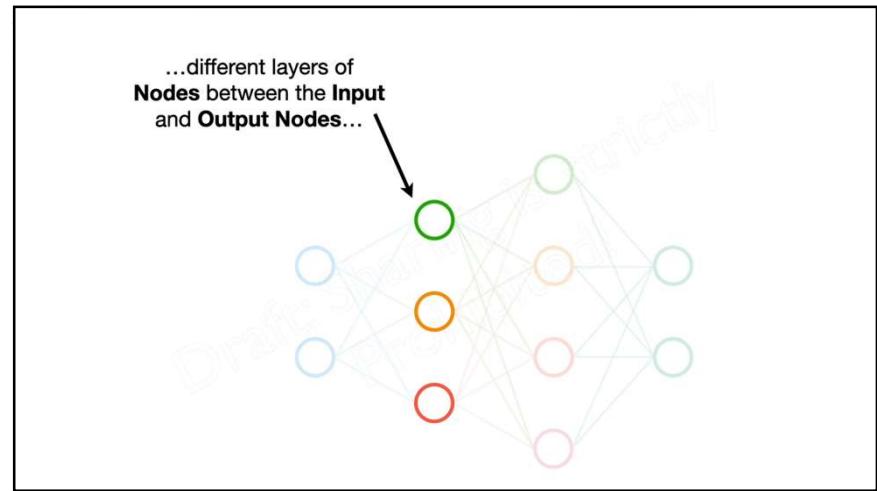
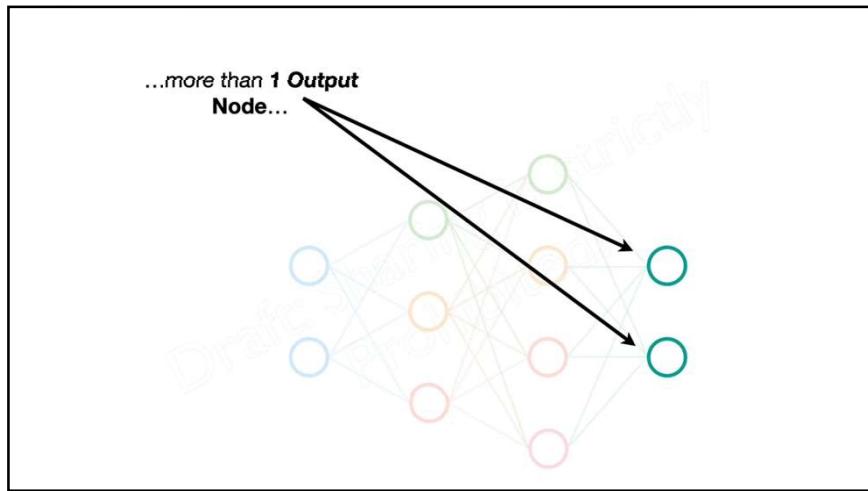
NOTE: This specific Neural Network is about as simple as they get.



It only has **1 Input Node** where we plug in the **Dosage**...







The layers of **Nodes** between the **Input** and **Output Nodes** are called **Hidden Layers**.



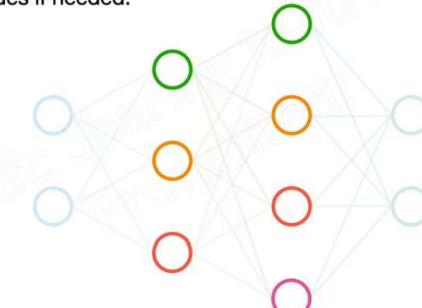
When you build a **Neural Network**, one of the first things you do is decide how many **Hidden Layers** you want, and how many **Nodes** go into each **Hidden Layer**.



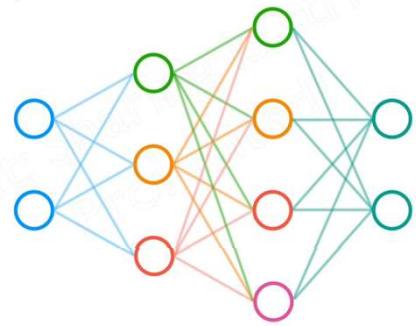
Although there are rules of thumb for making decisions about the **Hidden Layers**...



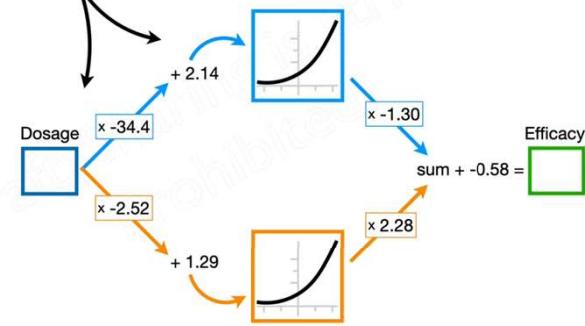
...you essentially make a guess and see how well the **Neural Network** performs, adding more layers and nodes if needed.



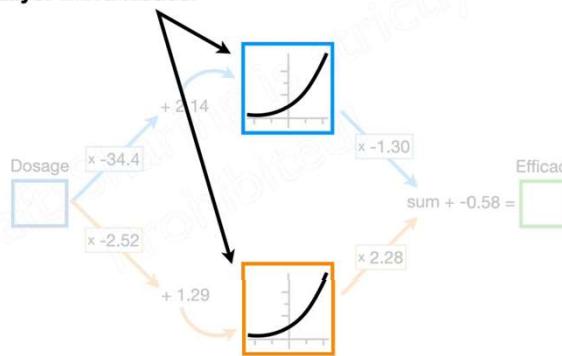
Now, even though this **Neural Network** looks fancy, it is still made from the same parts...



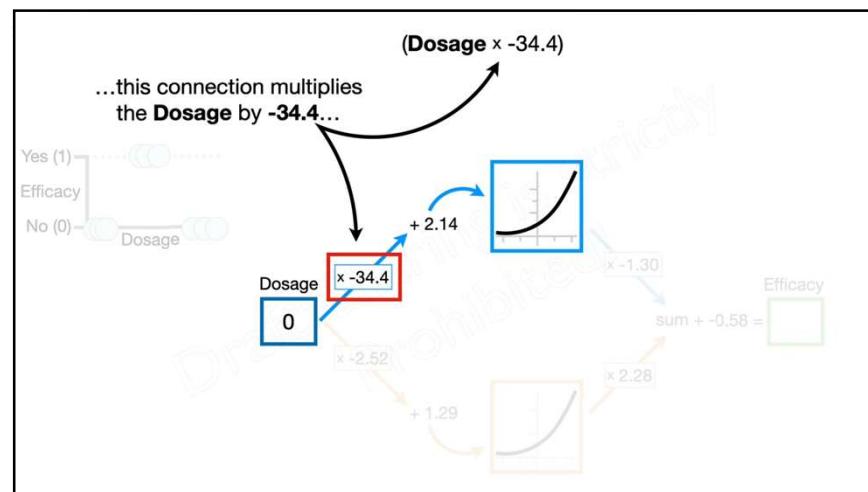
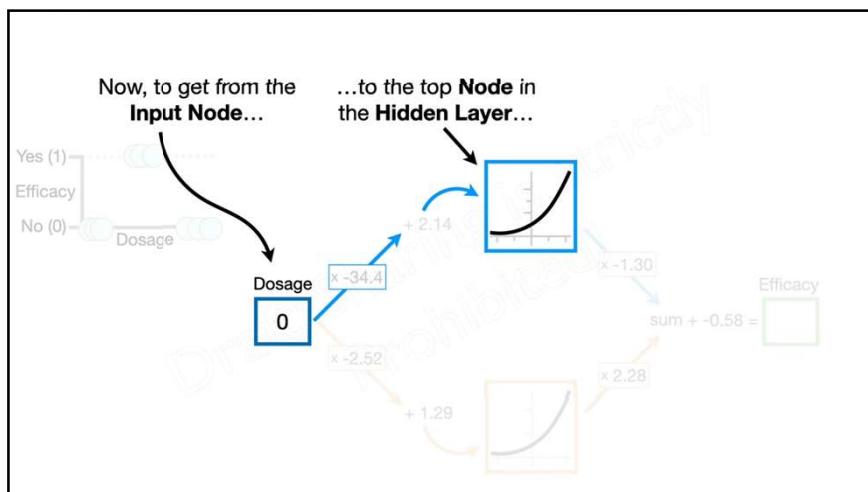
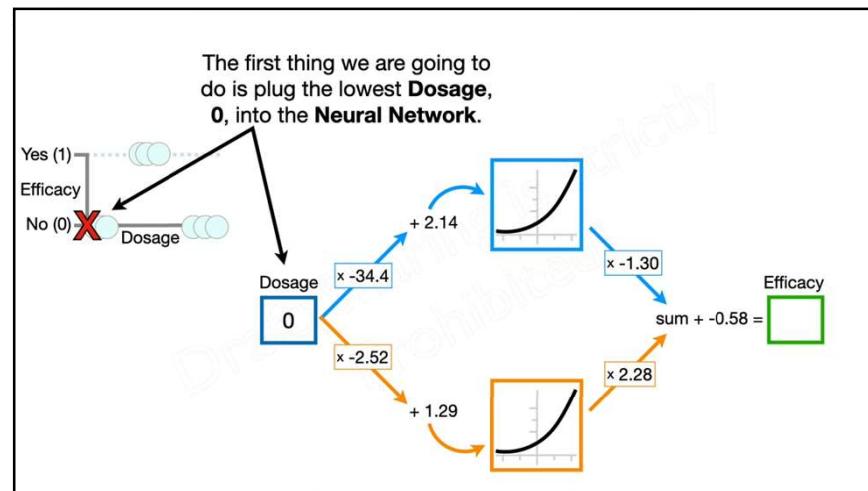
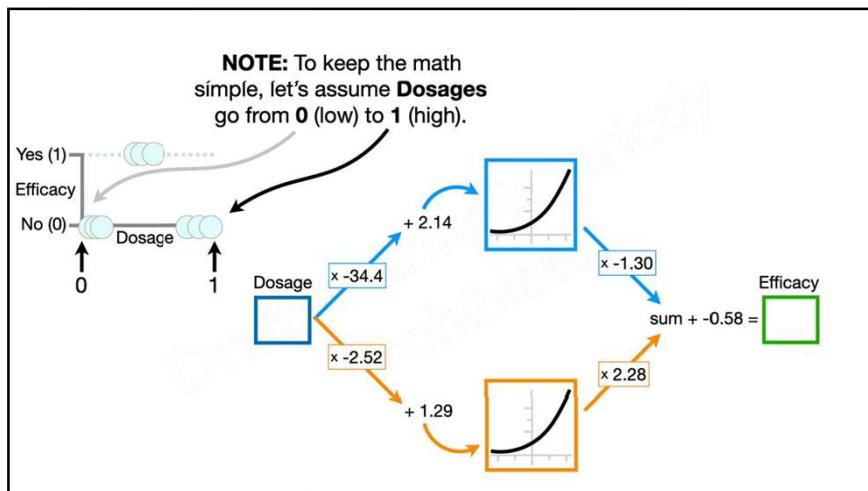
...used in this simple
Neural Network...

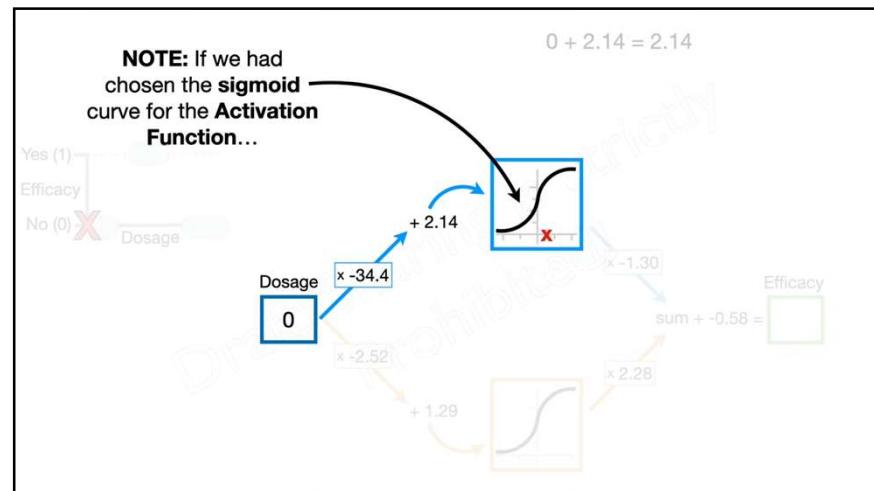
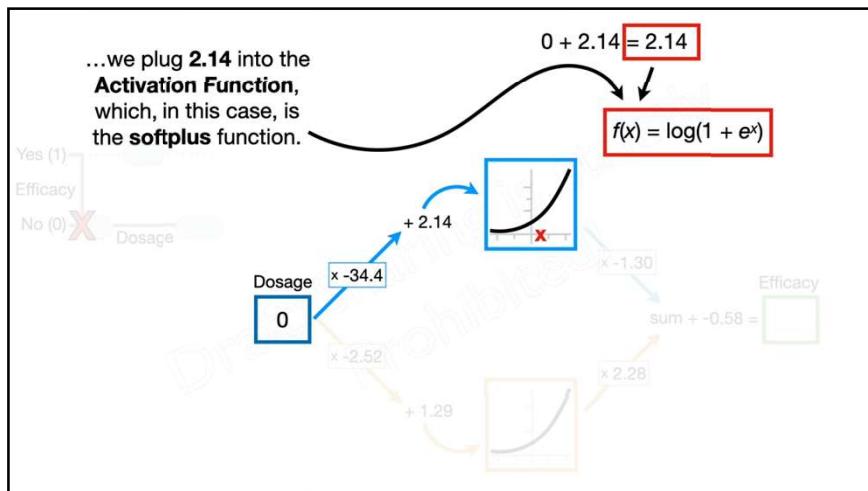
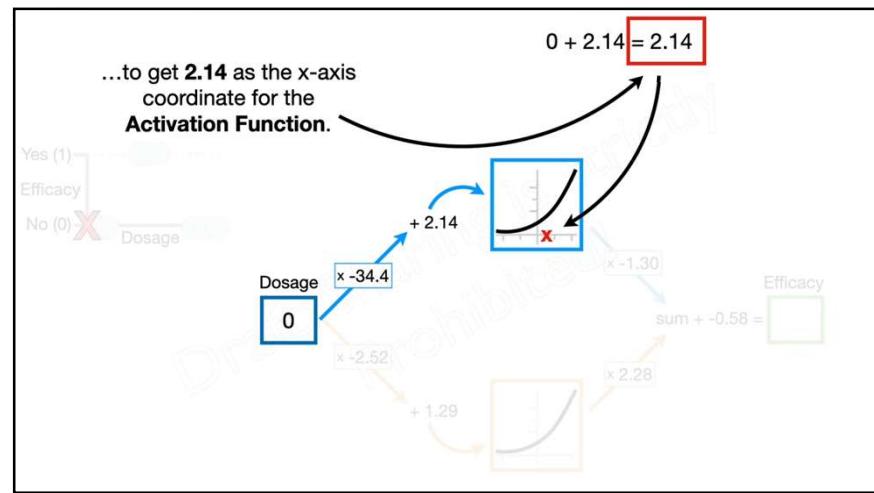
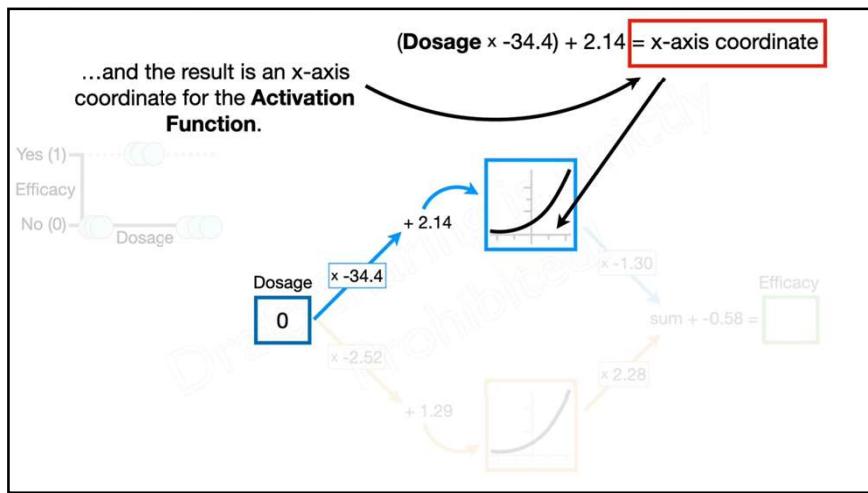


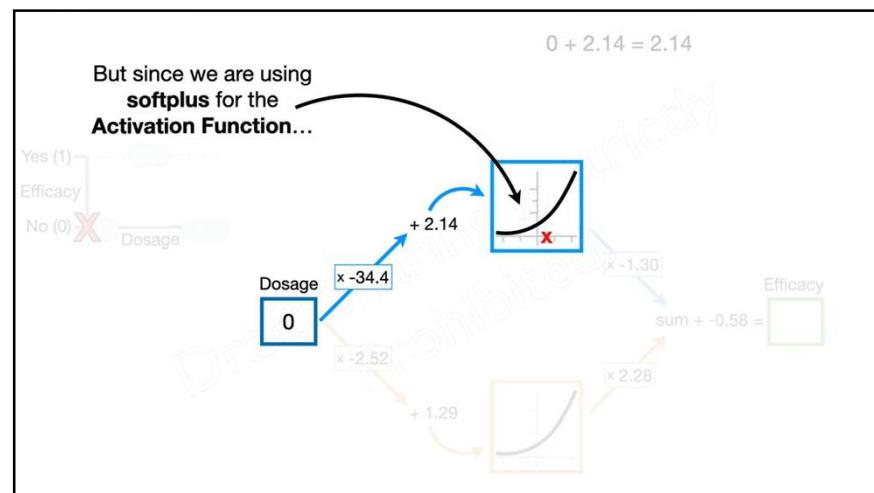
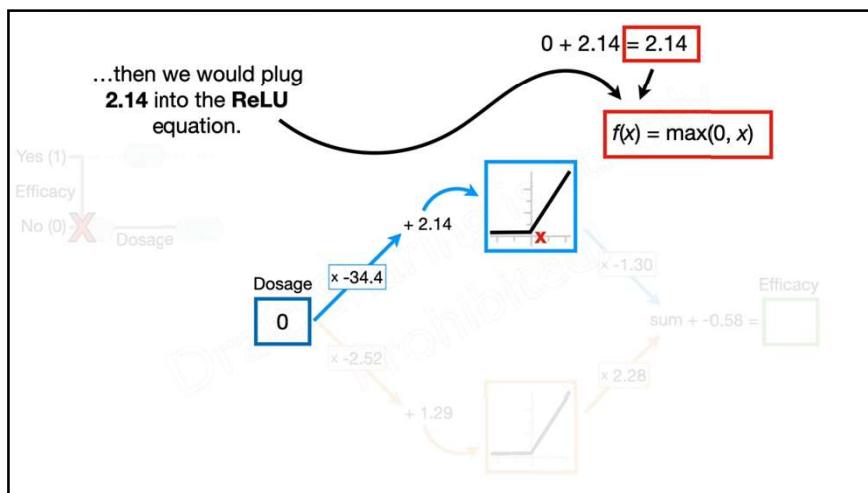
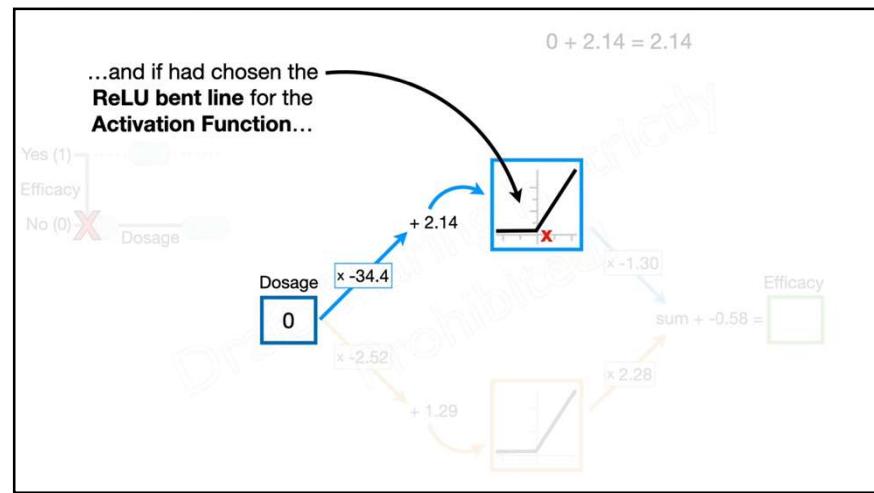
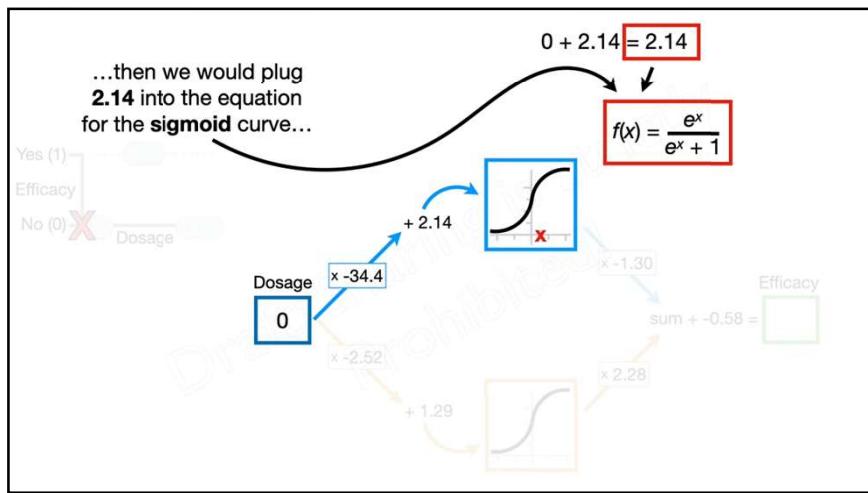
...which only has **1 Hidden Layer** with **2 Nodes**.

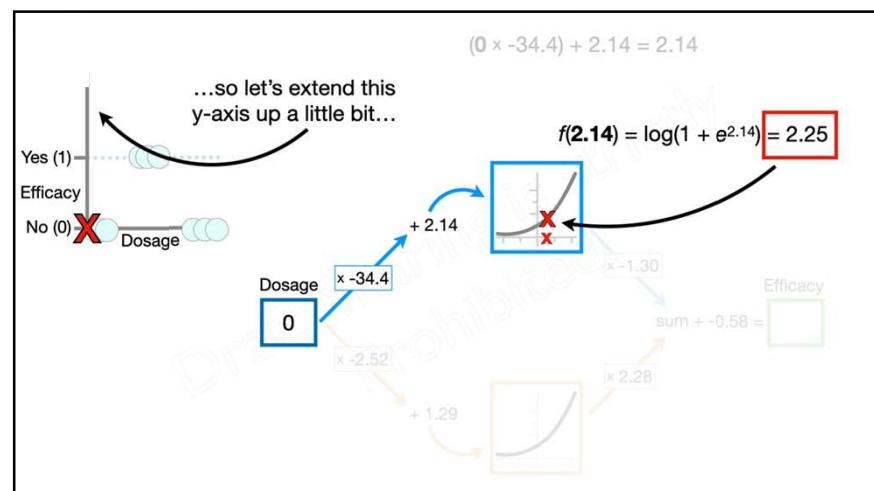
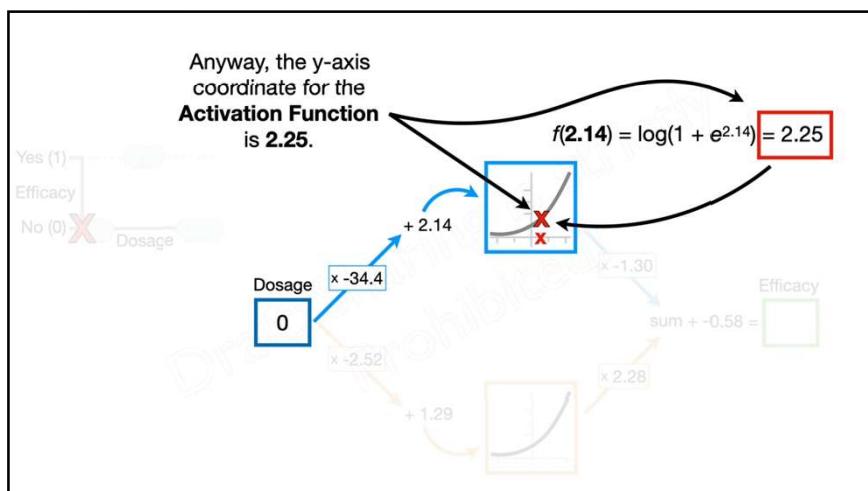
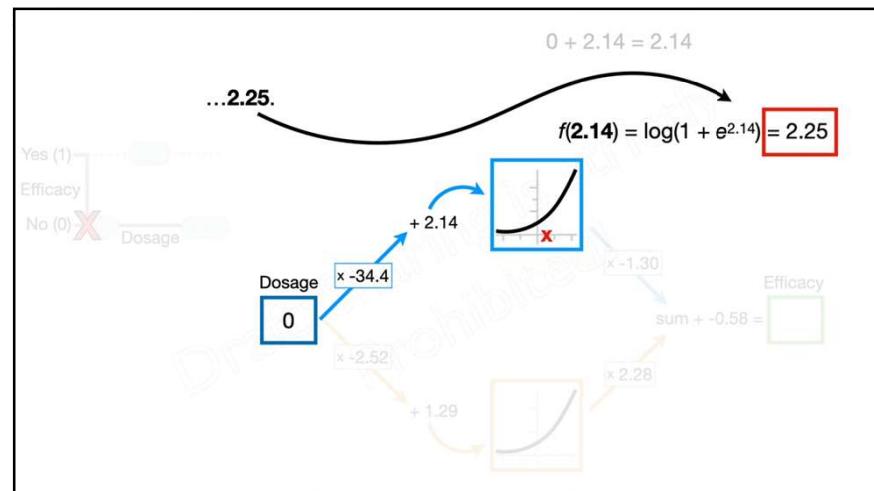
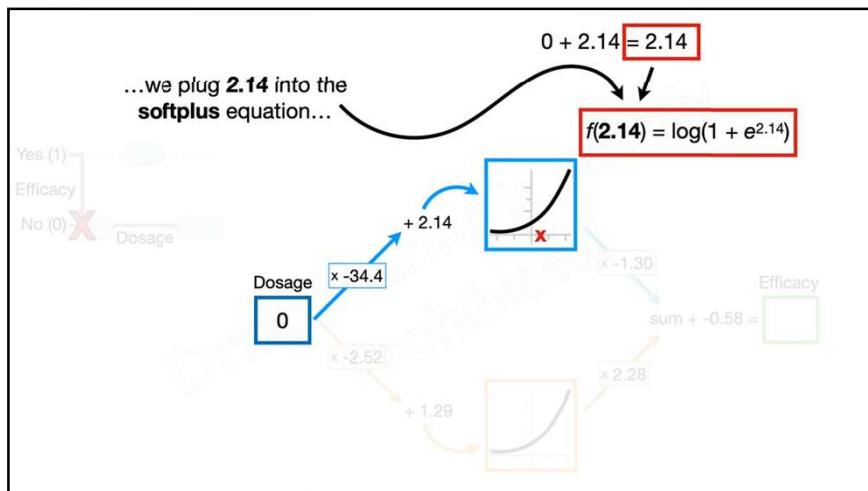


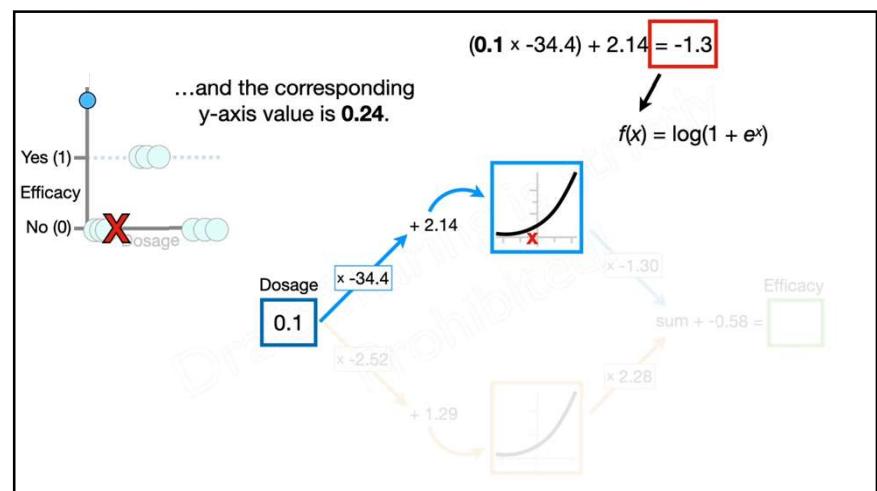
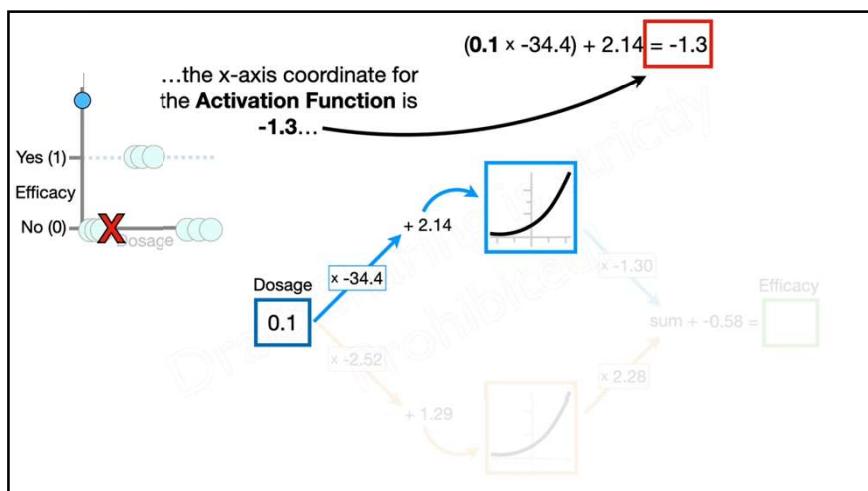
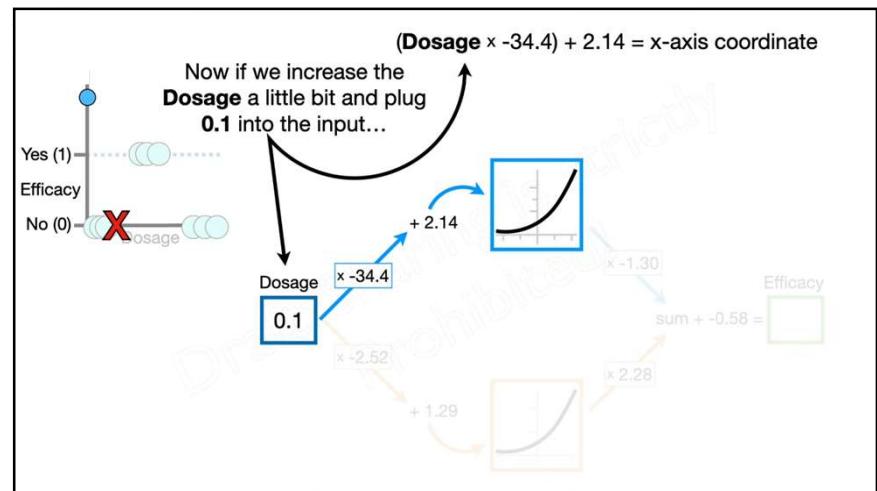
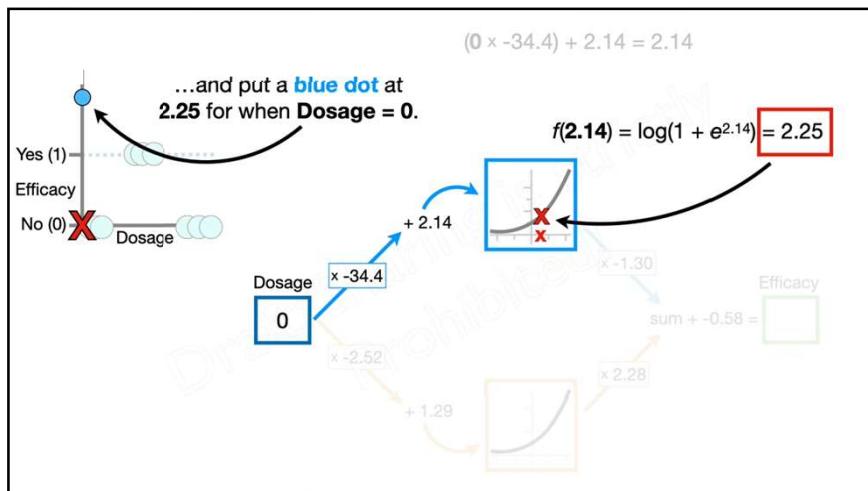
So let's learn how this **Neural Network** creates new shapes from the **curved or bent lines** in the **Hidden Layer**...

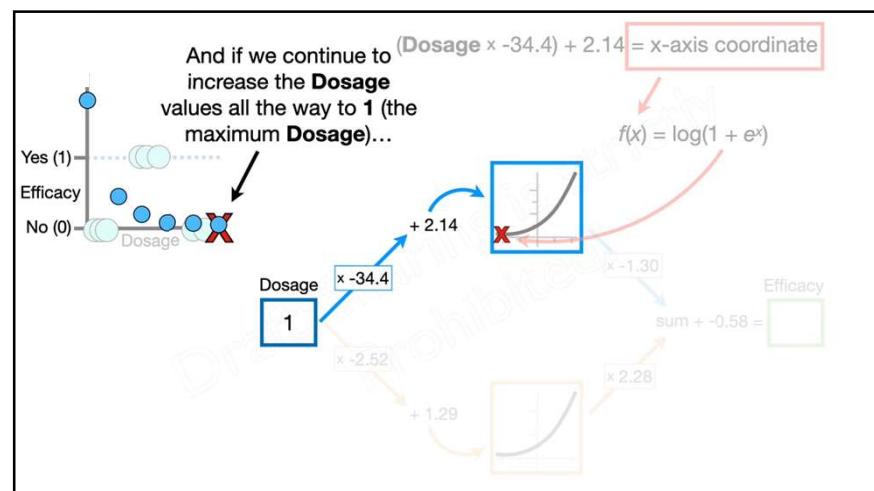
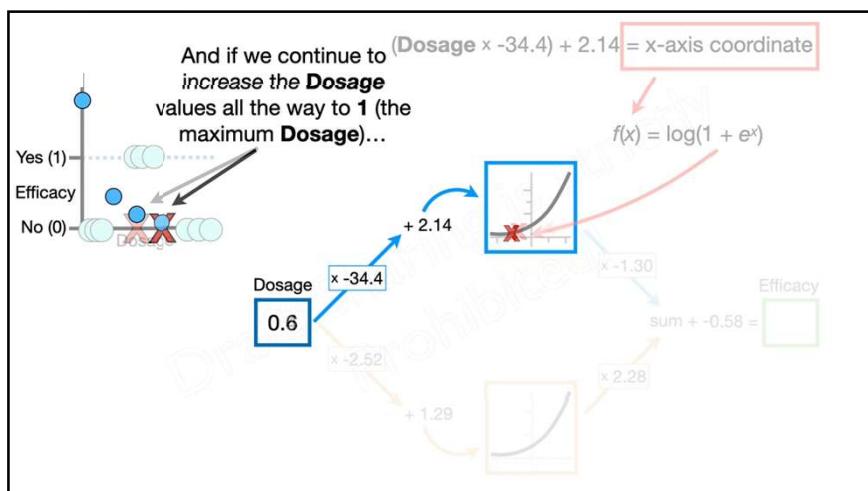
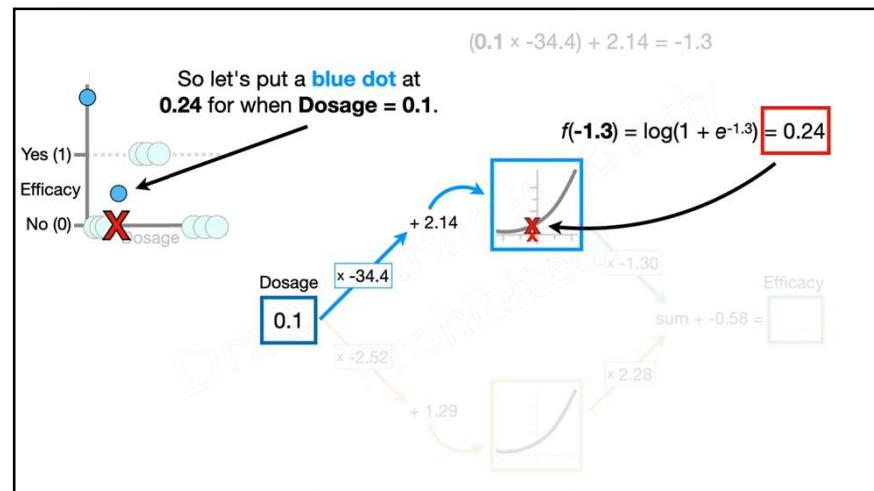
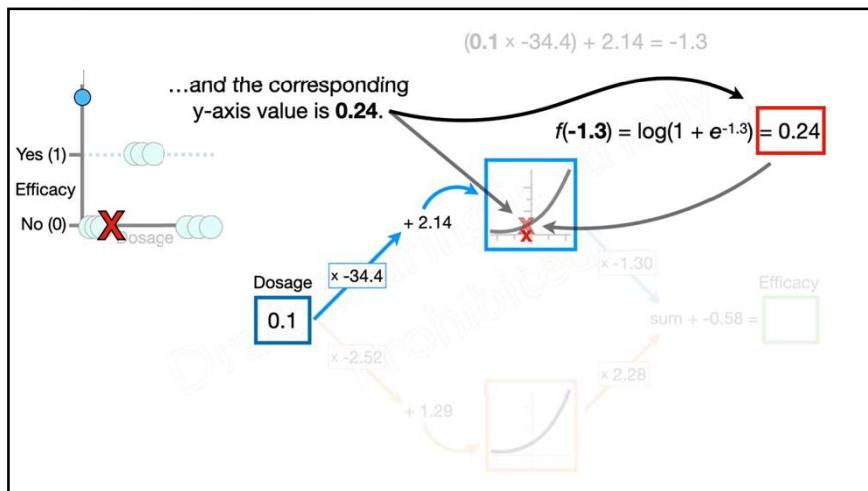


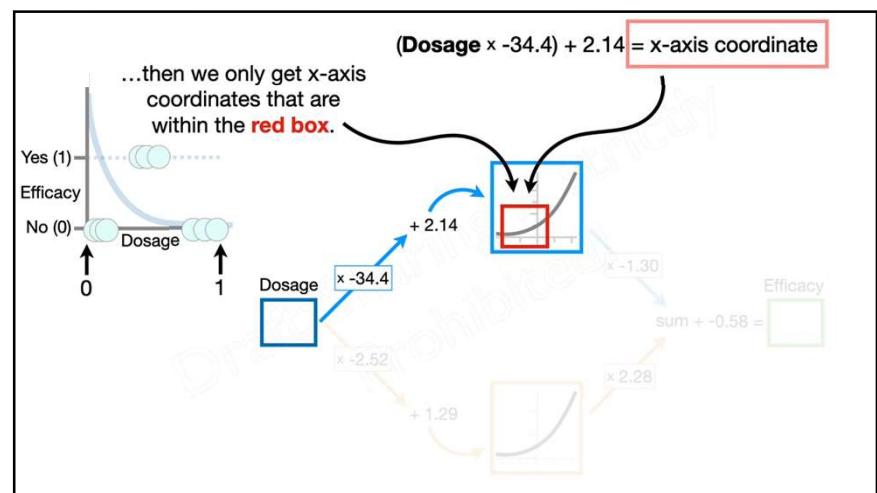
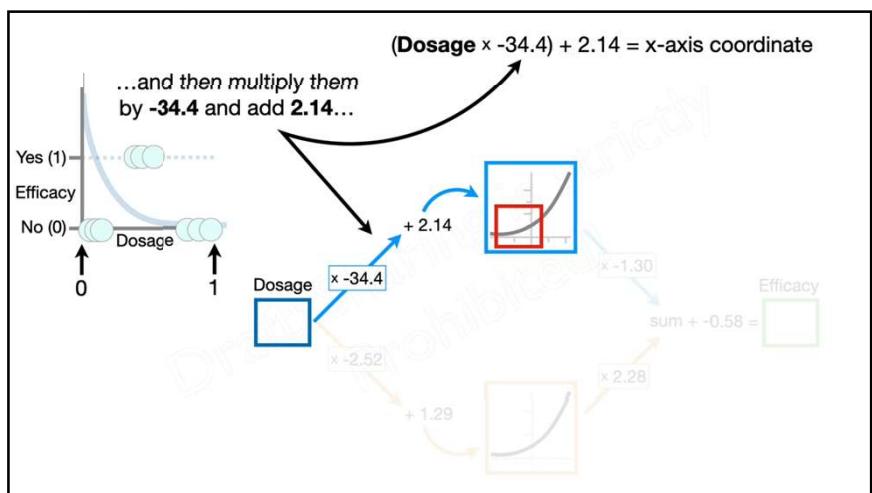
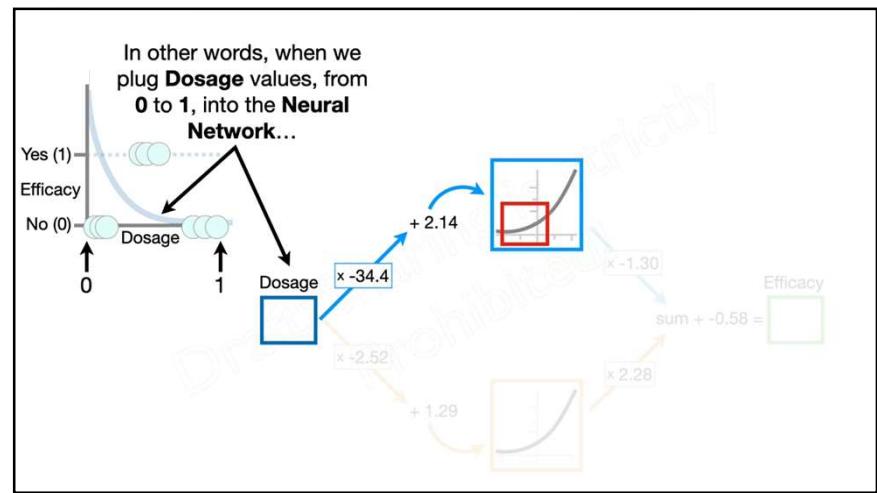
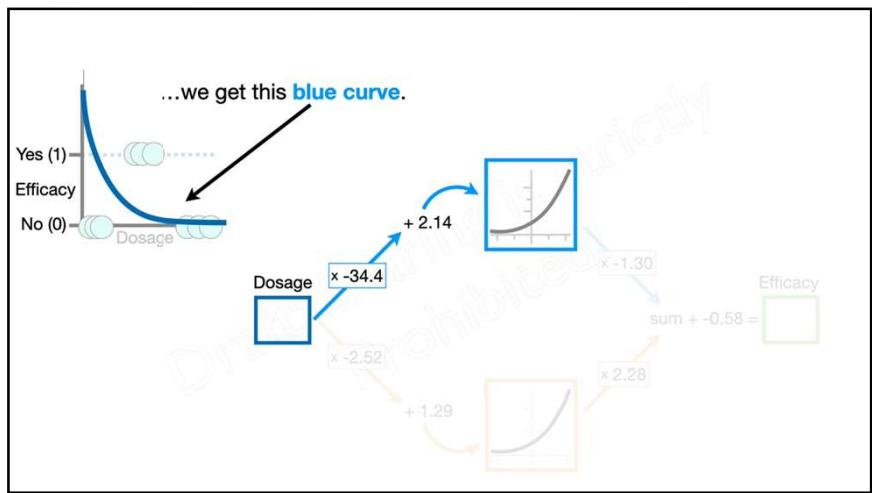


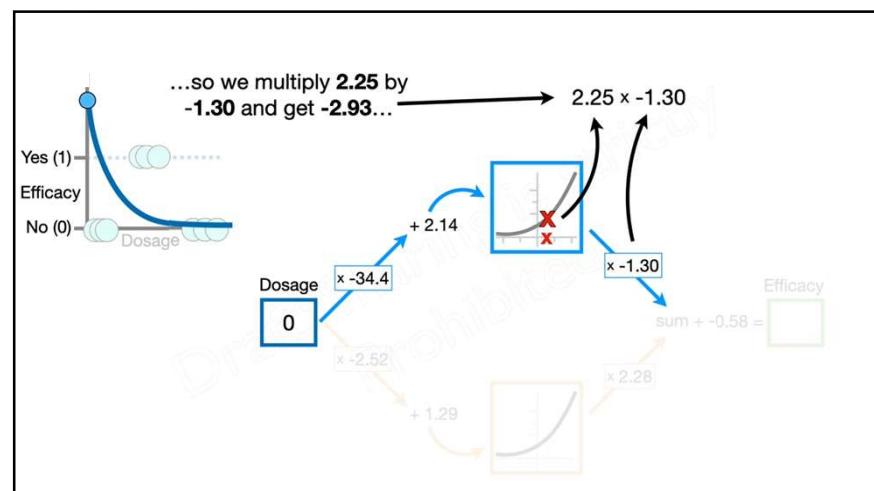
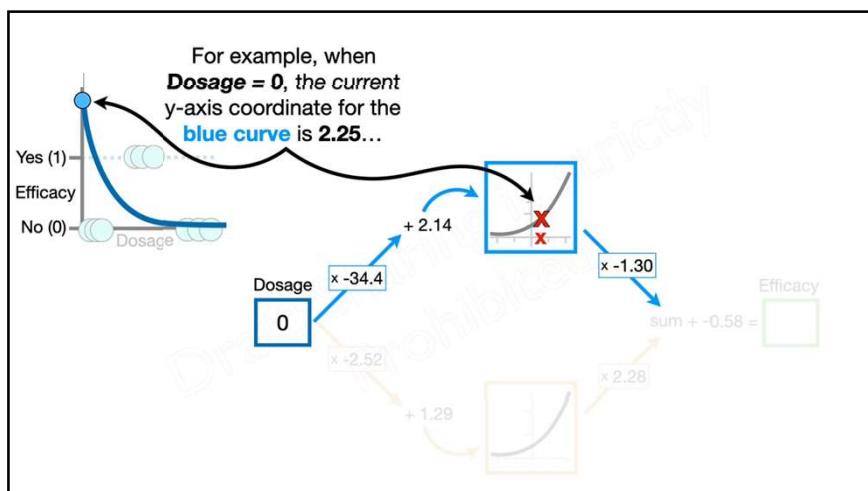
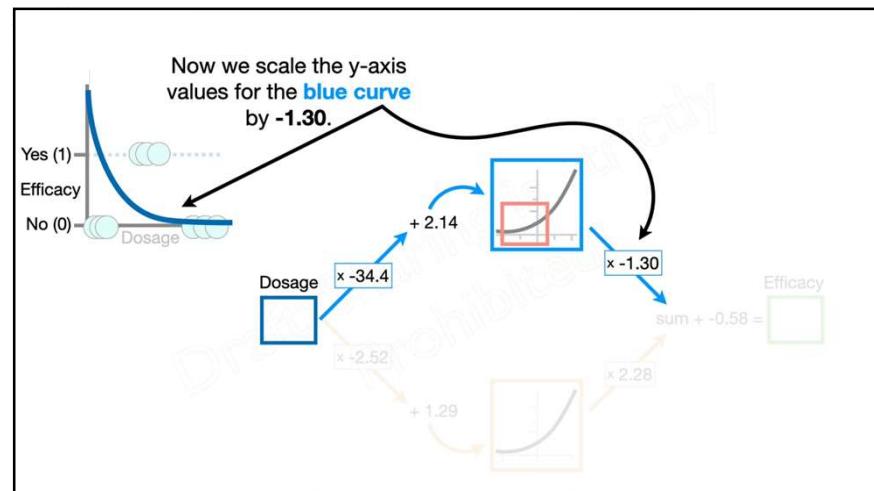
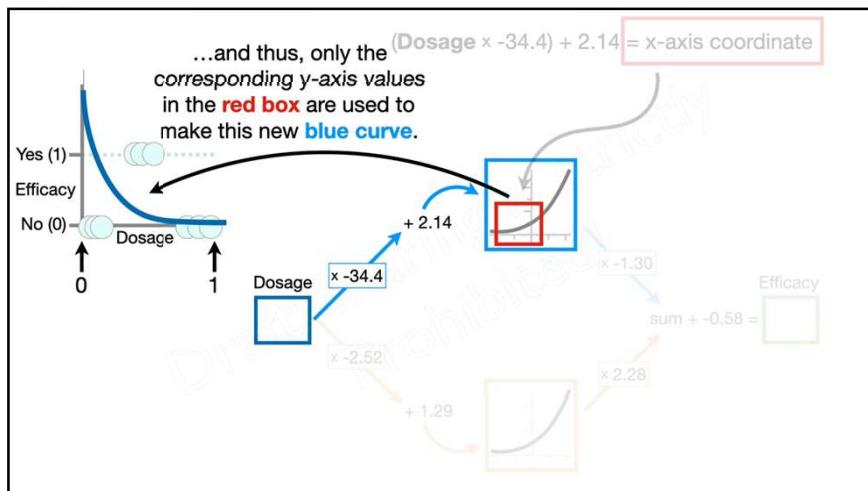


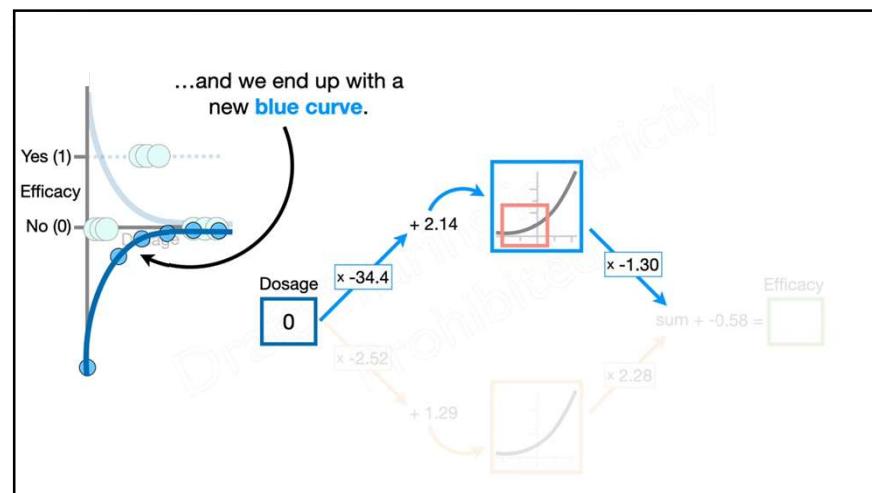
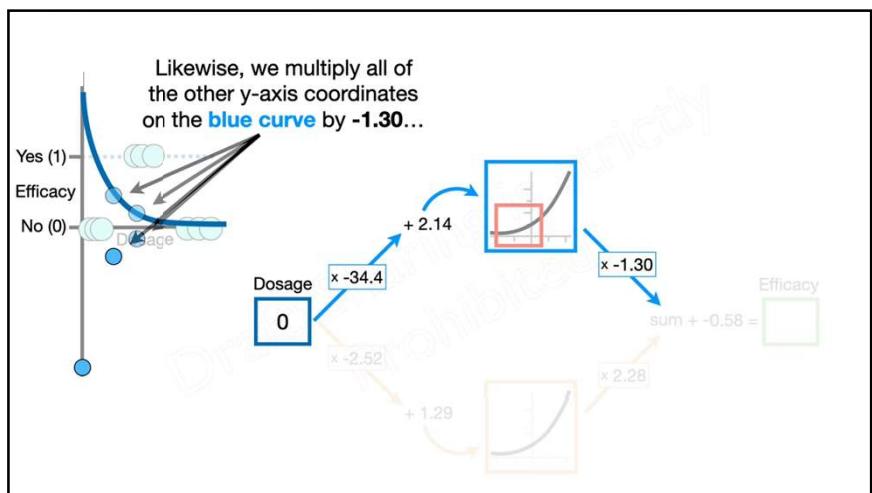
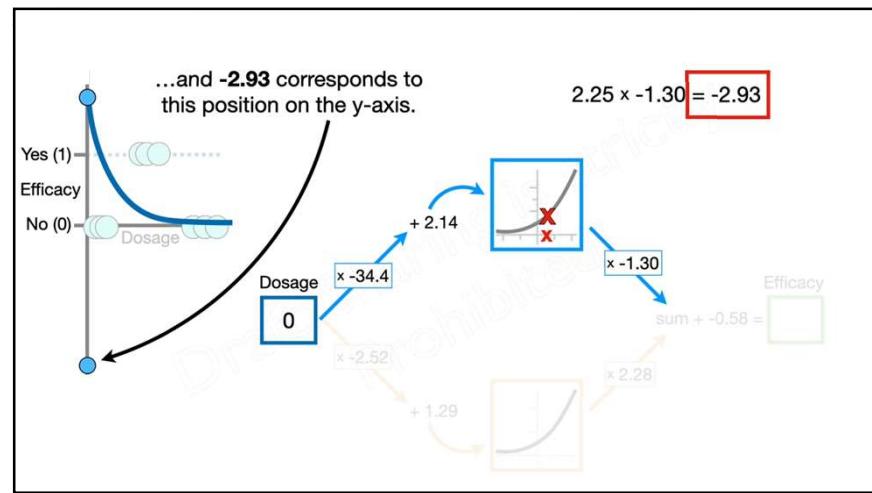
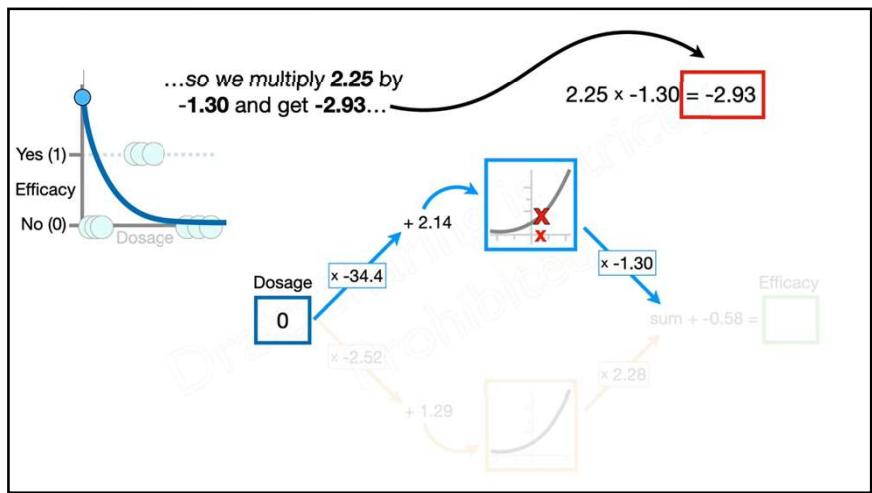


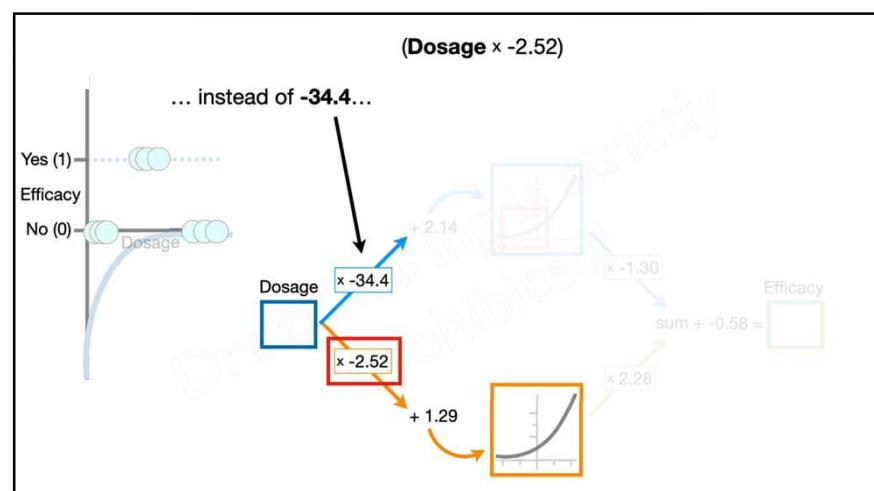
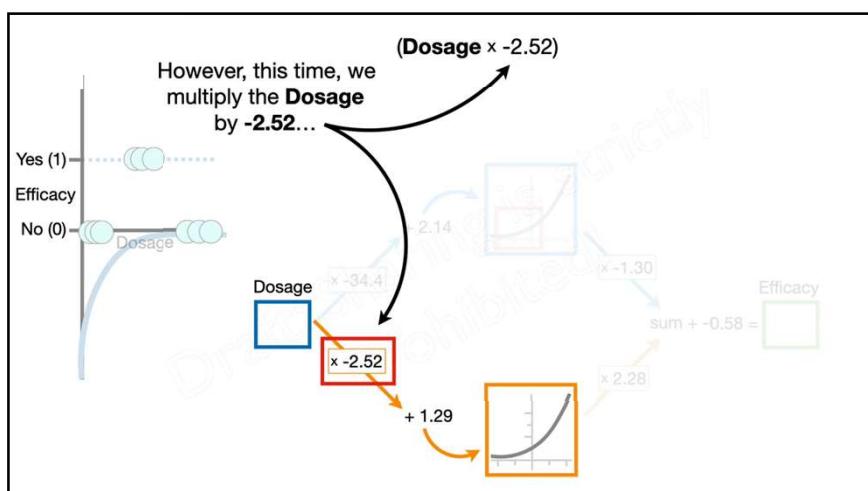
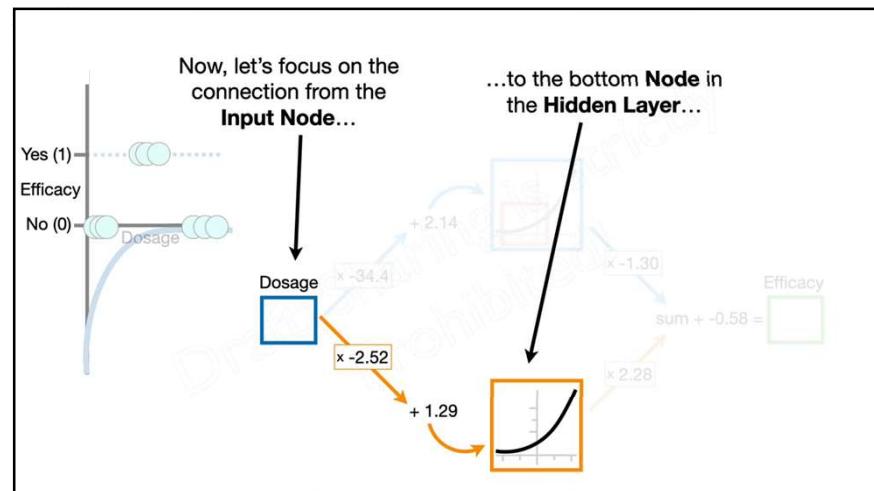
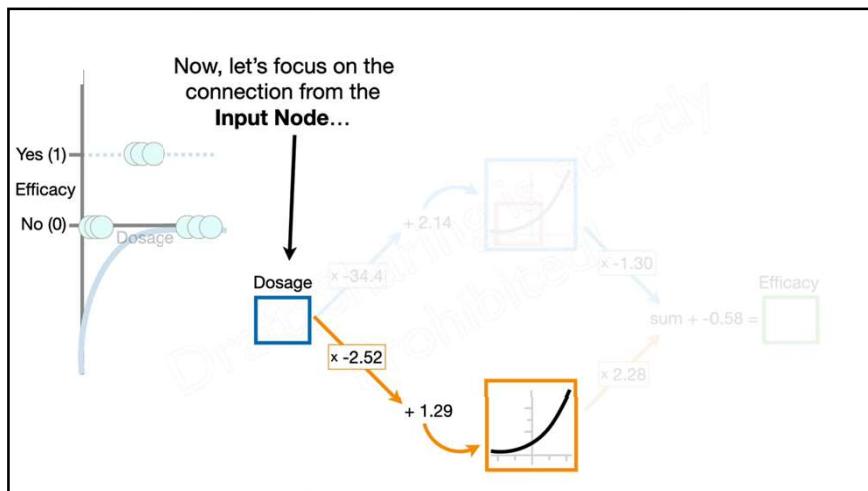


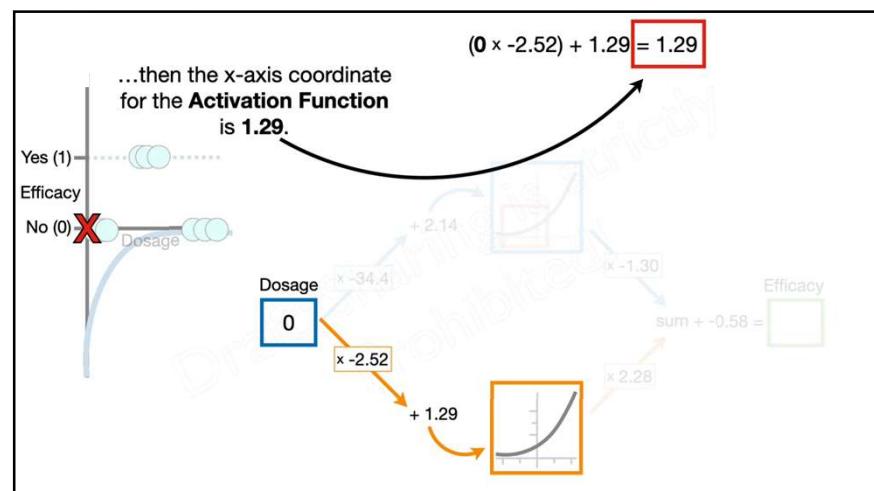
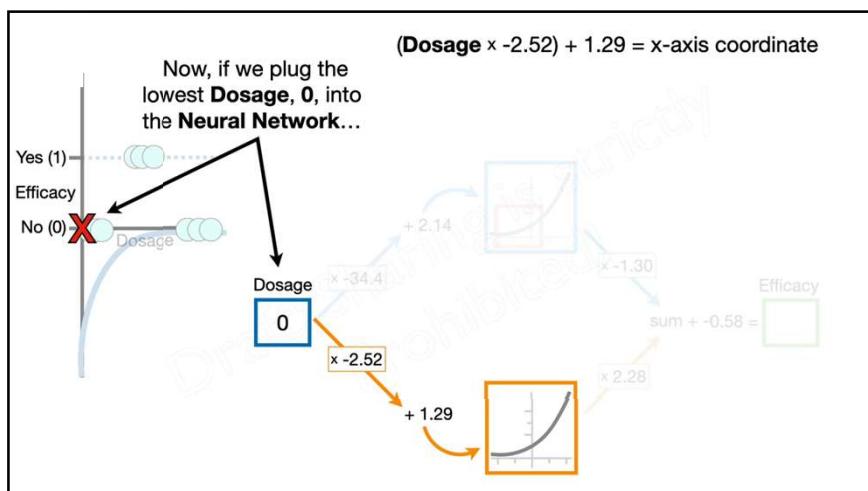
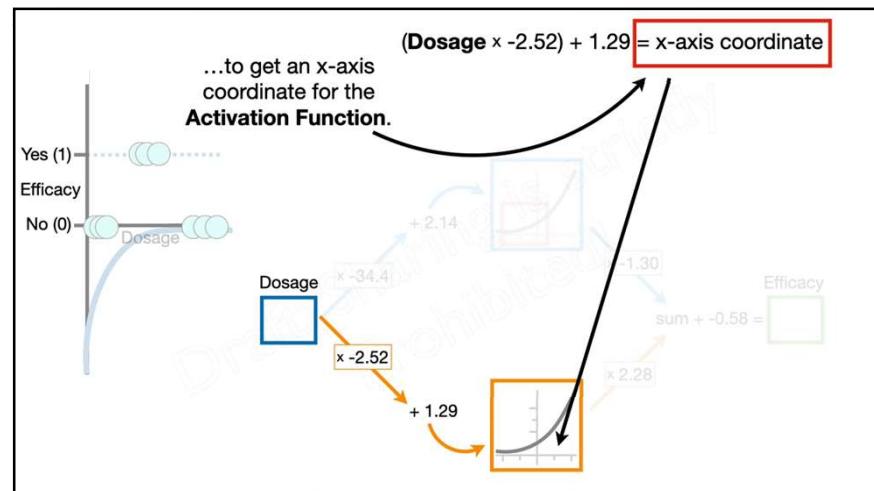
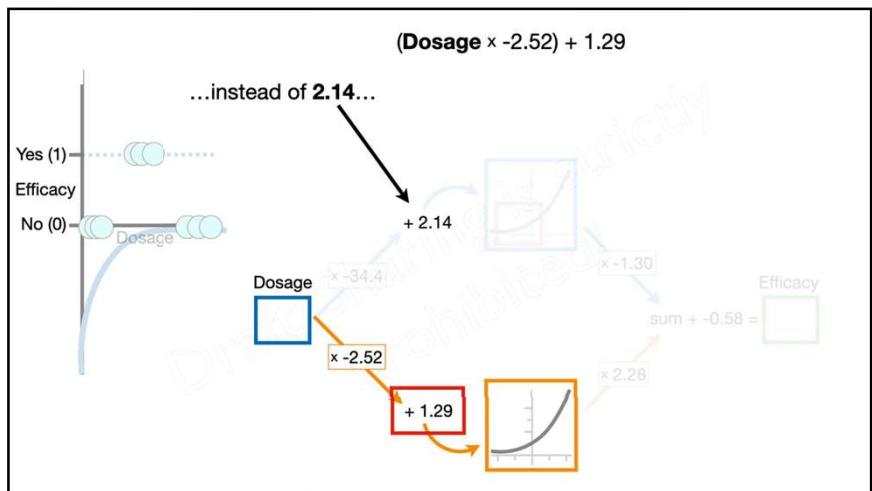


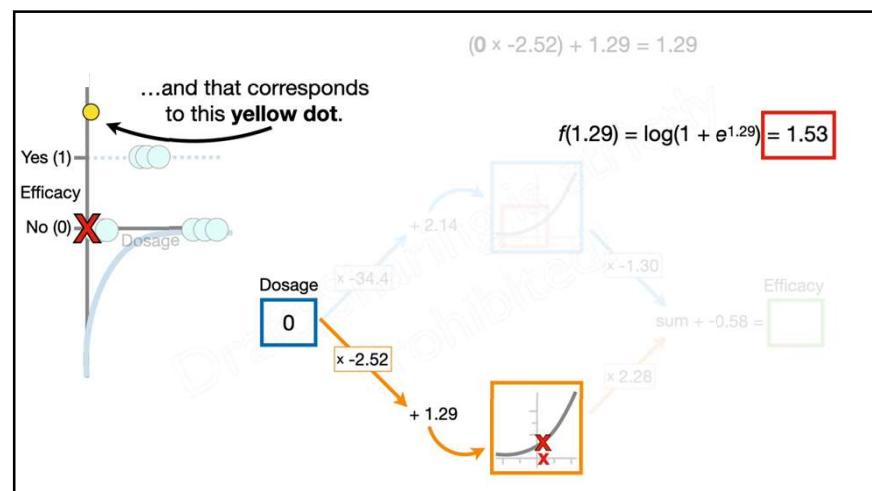
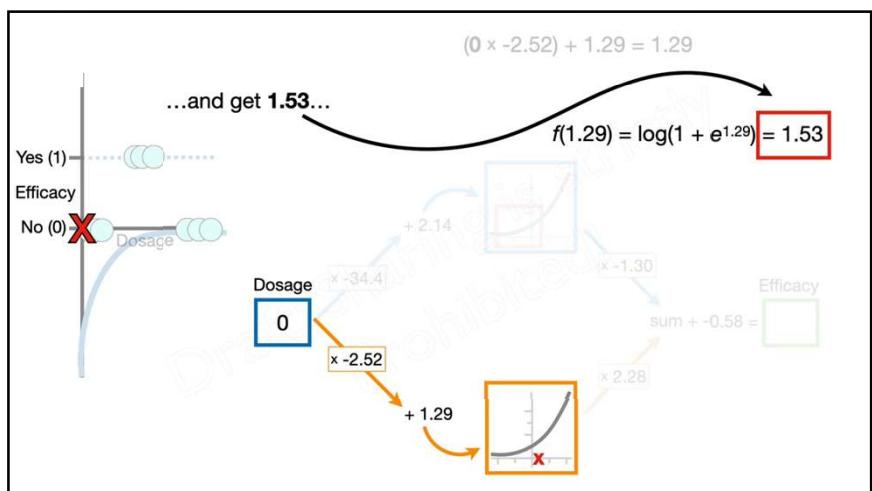
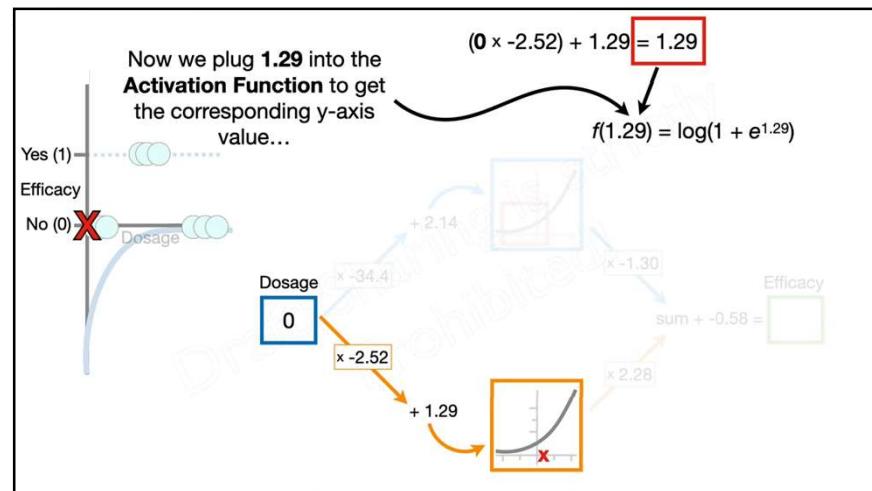
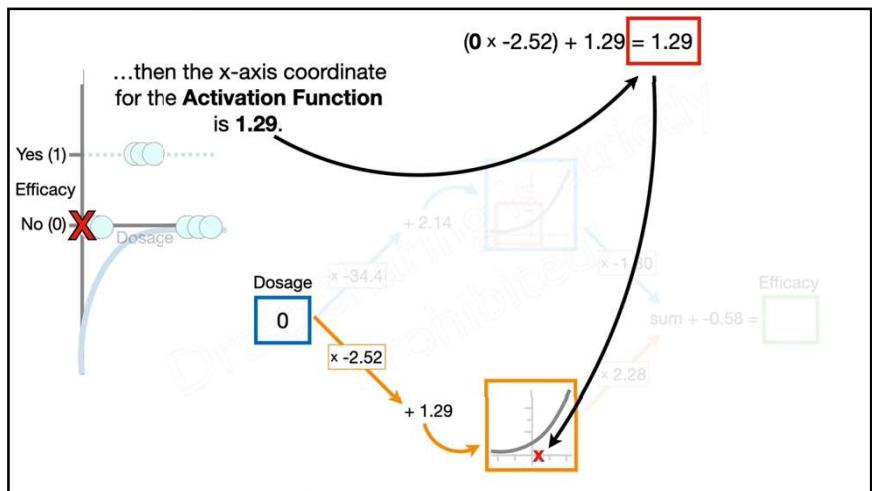


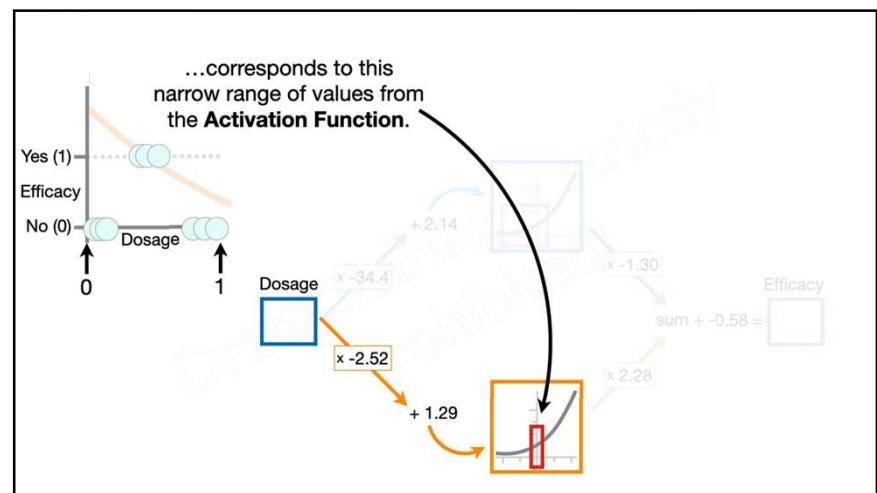
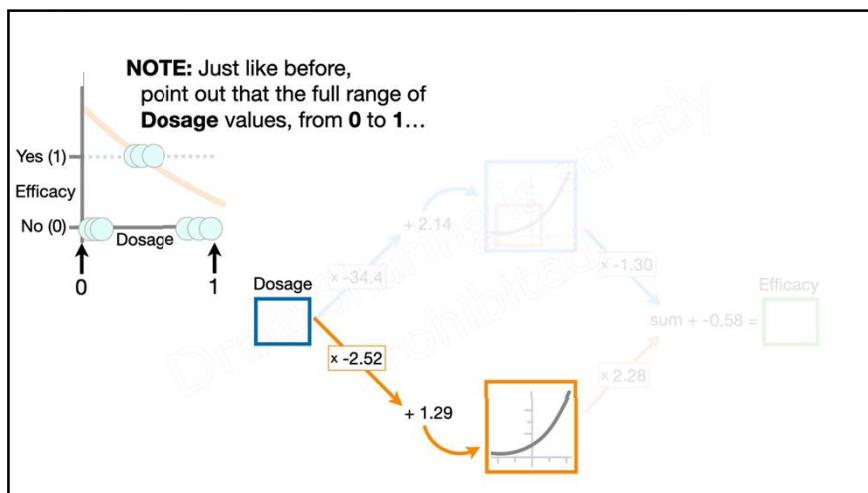
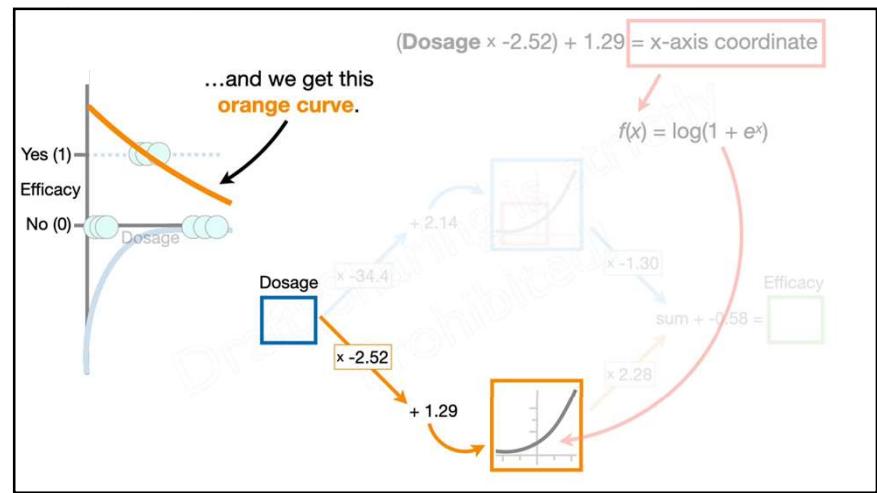
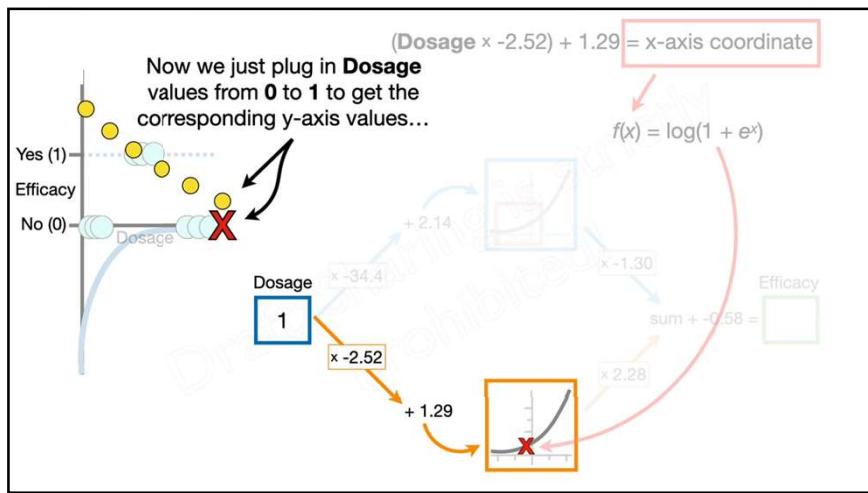




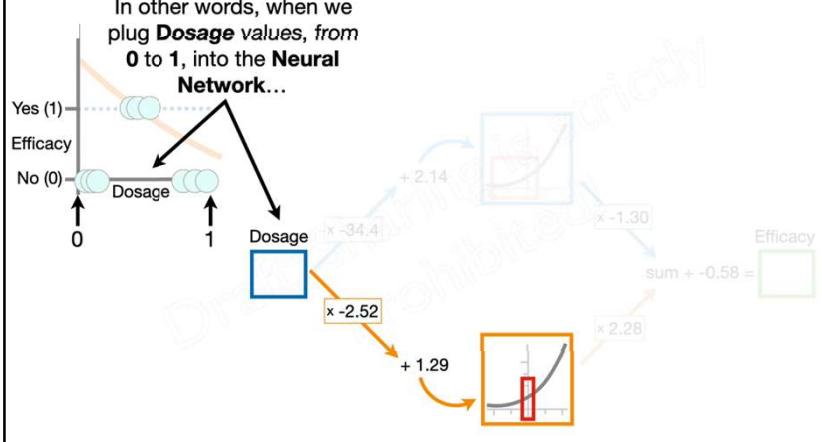






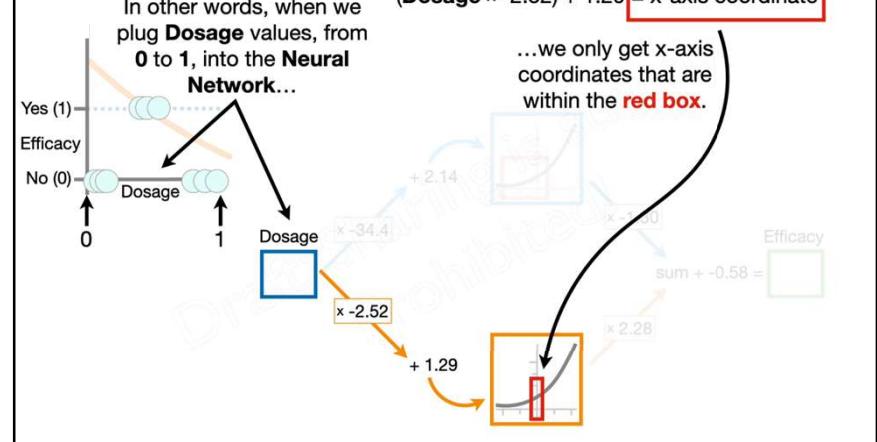


In other words, when we plug **Dosage** values, from 0 to 1, into the **Neural Network**...

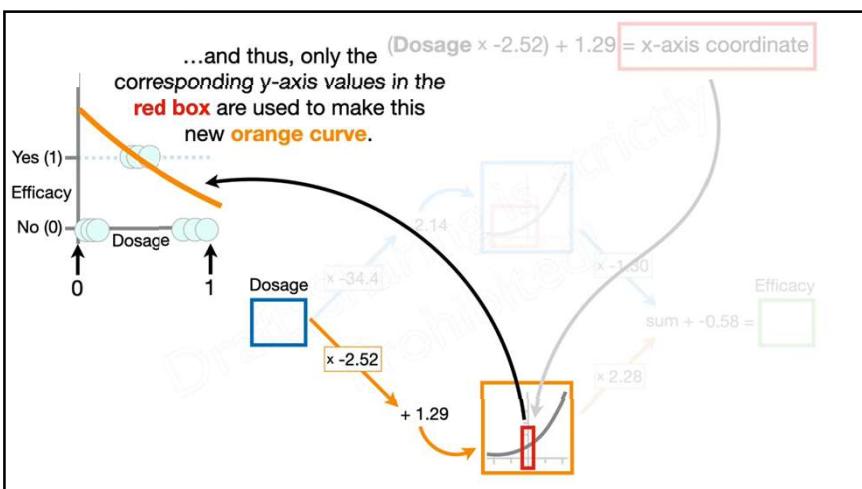


In other words, when we plug **Dosage** values, from 0 to 1, into the **Neural Network**...

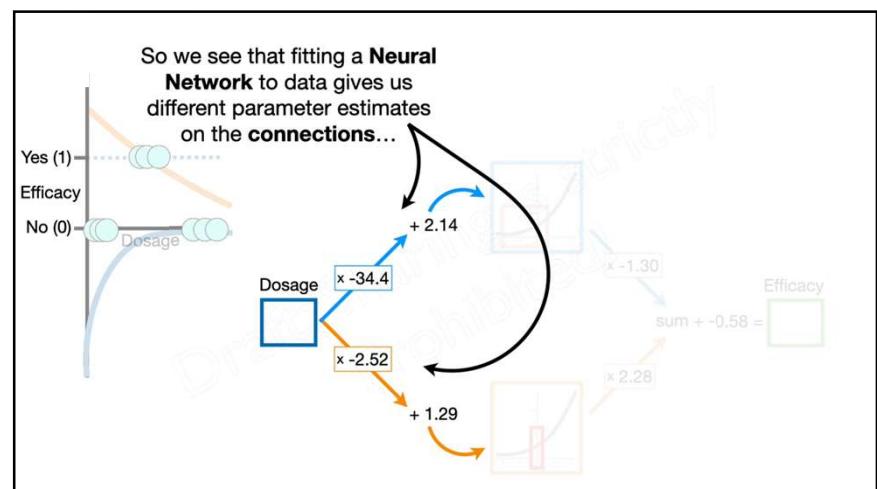
...we only get x-axis coordinates that are within the **red box**.)

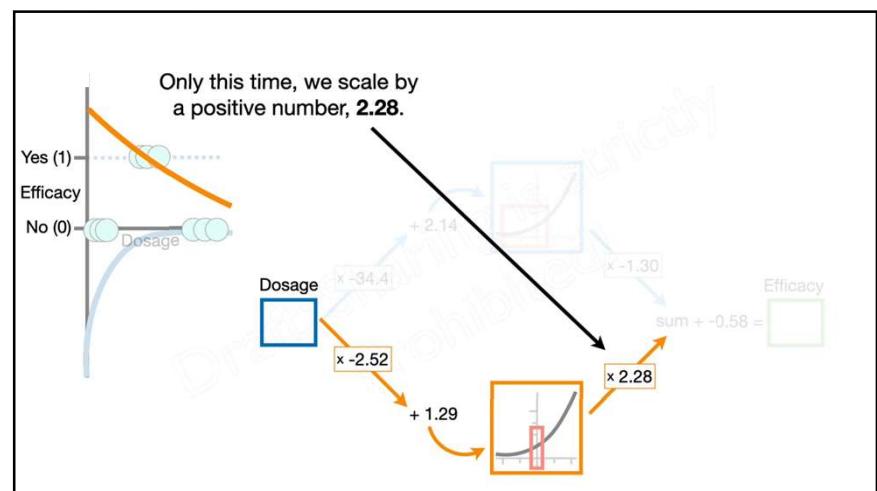
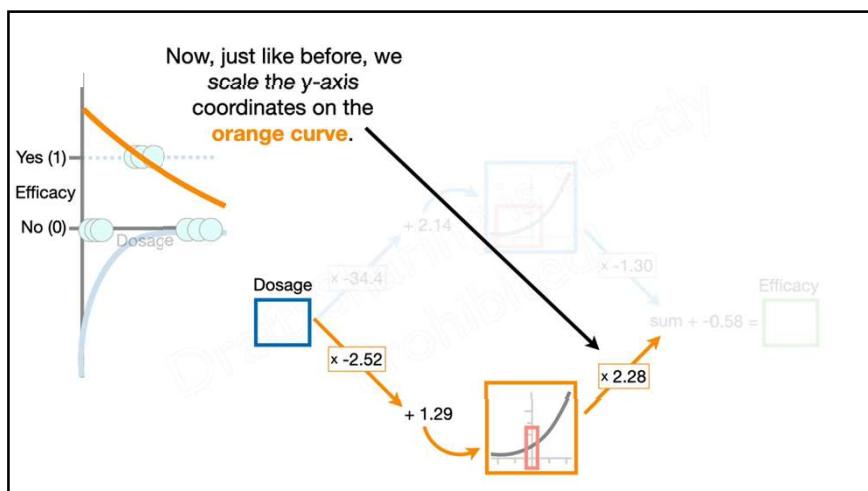
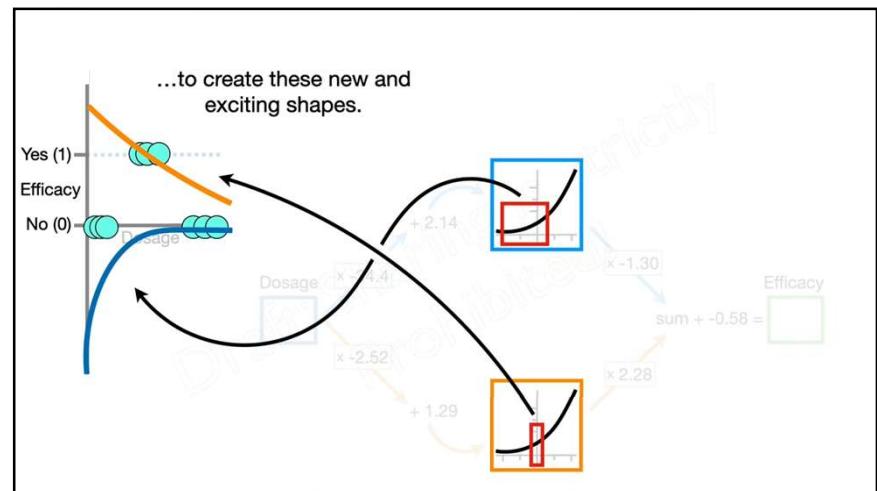
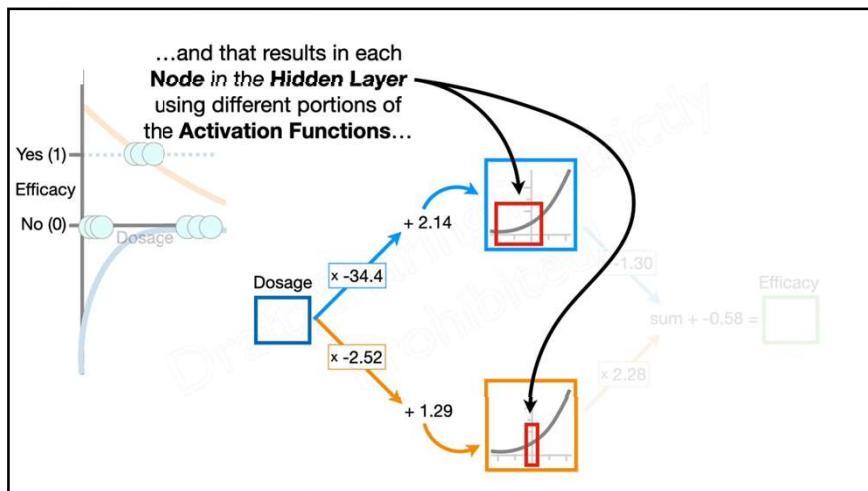


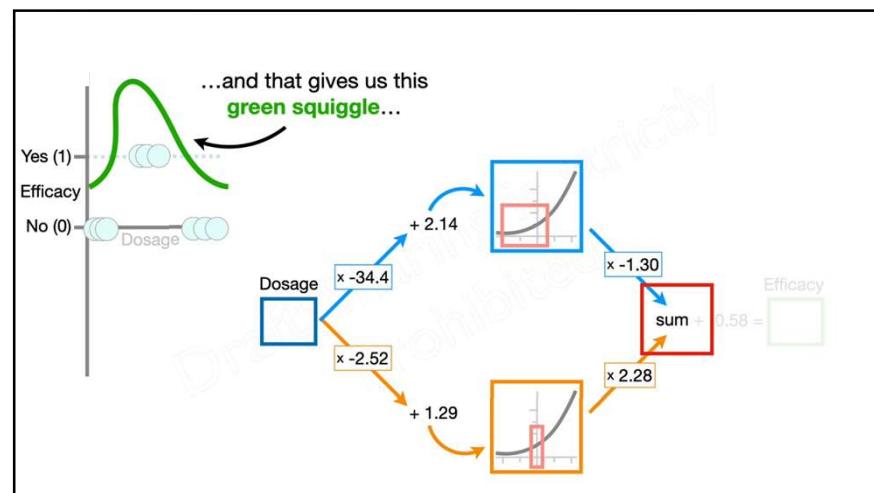
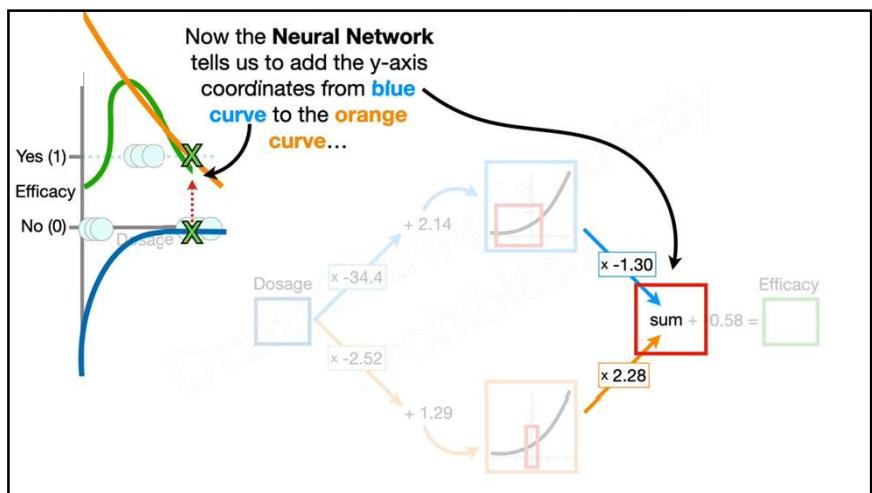
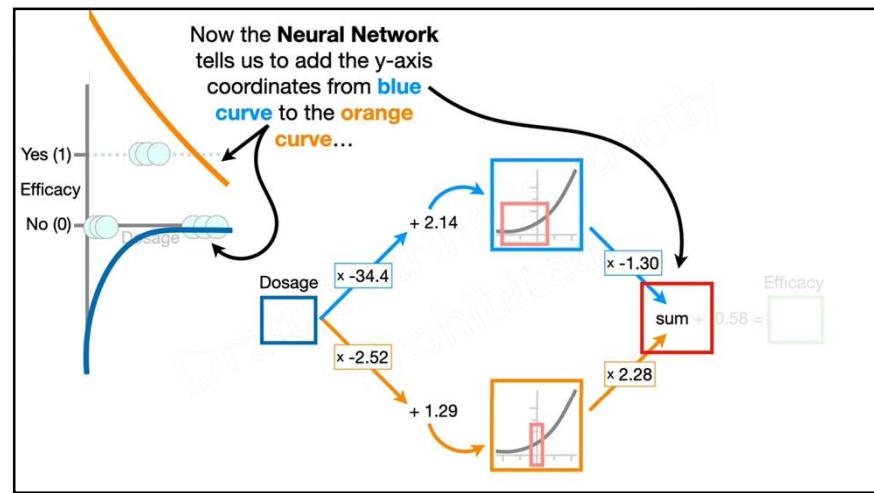
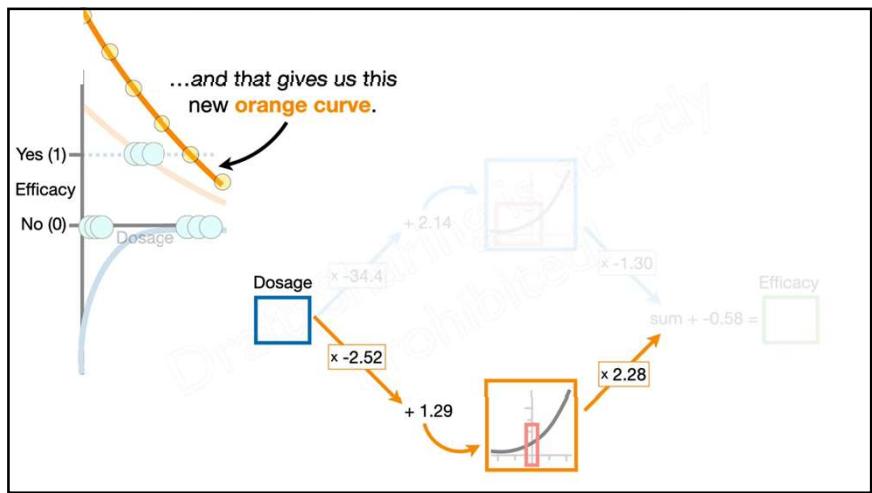
...and thus, only the corresponding y-axis values in the red box are used to make this new orange curve.

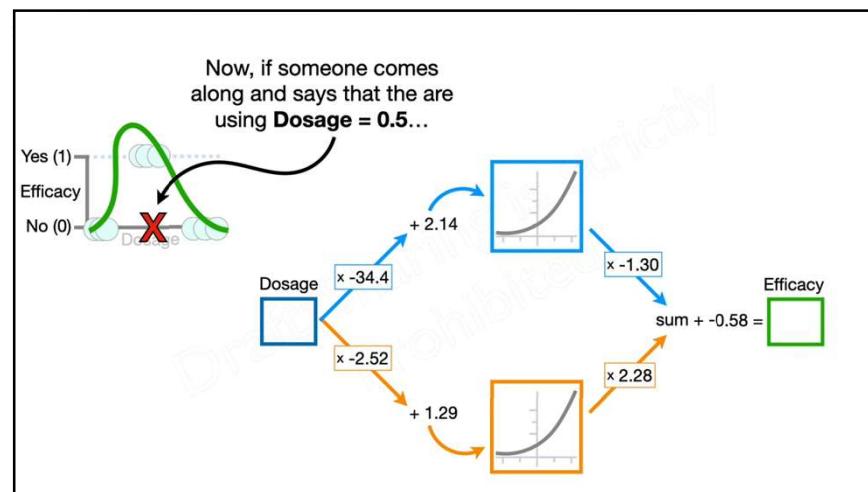
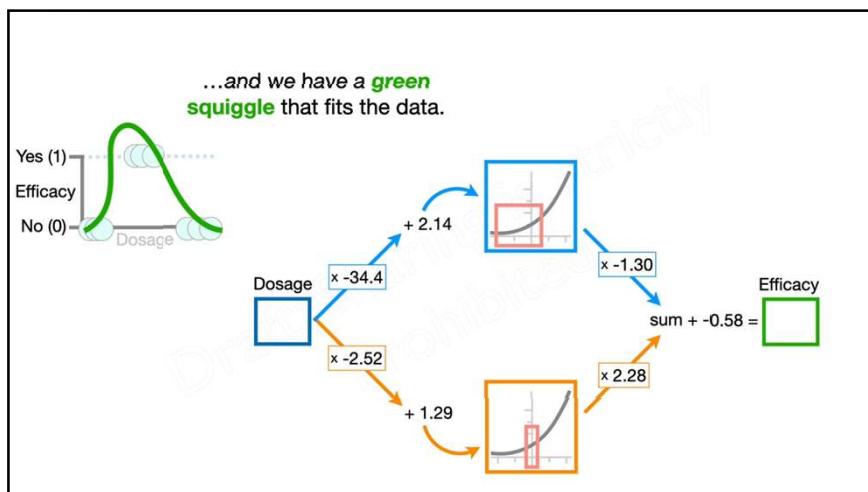
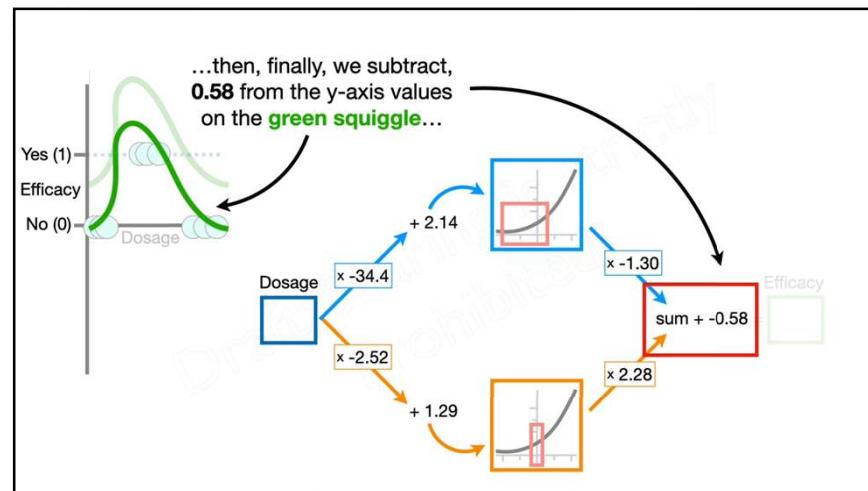
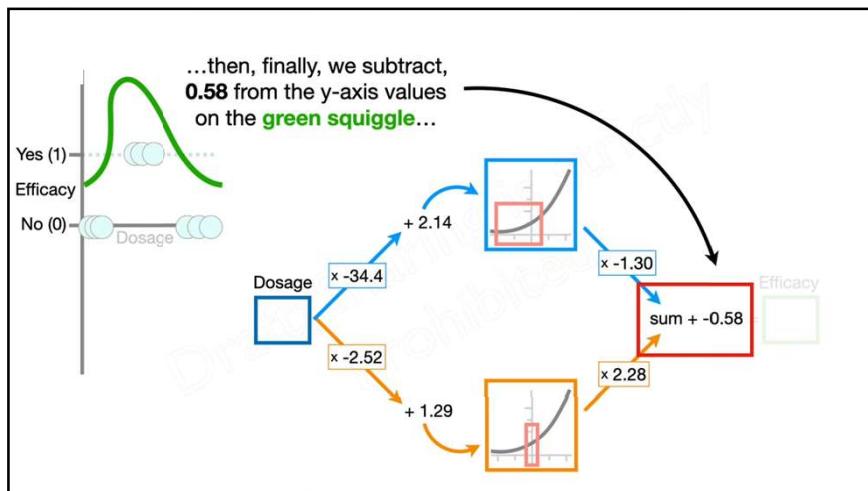


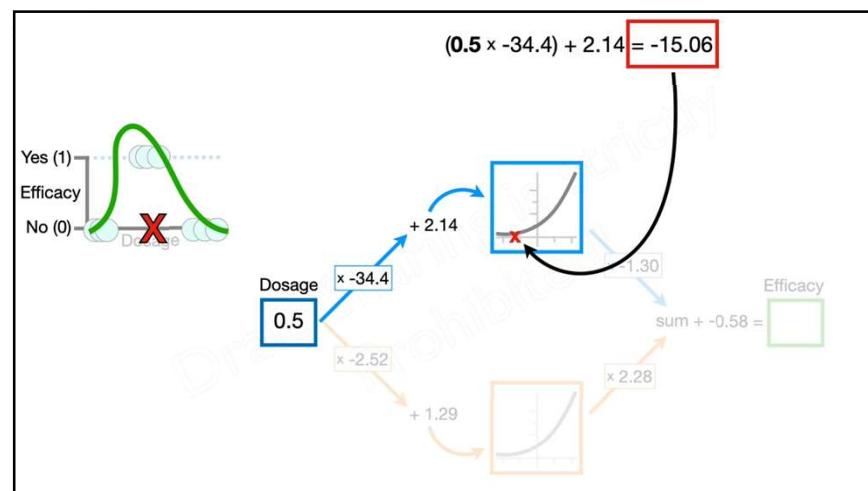
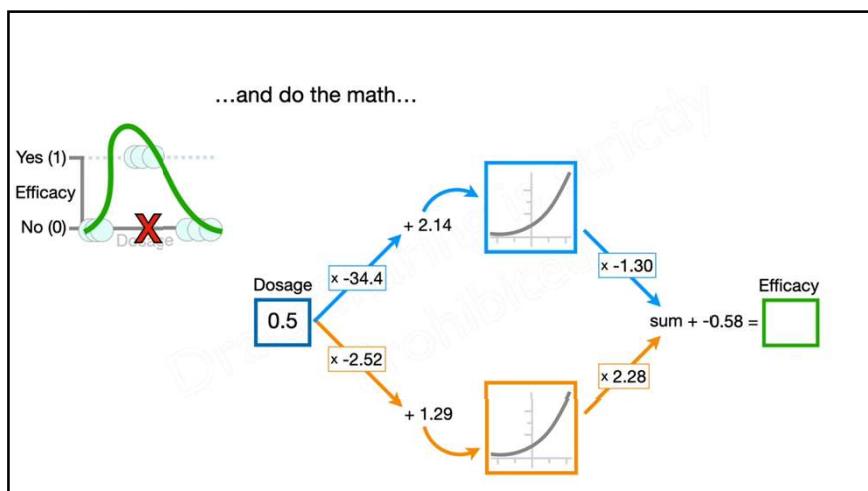
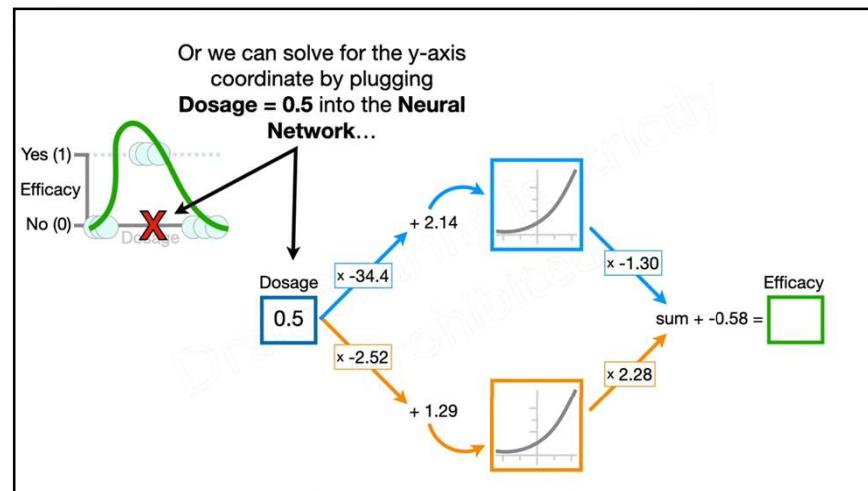
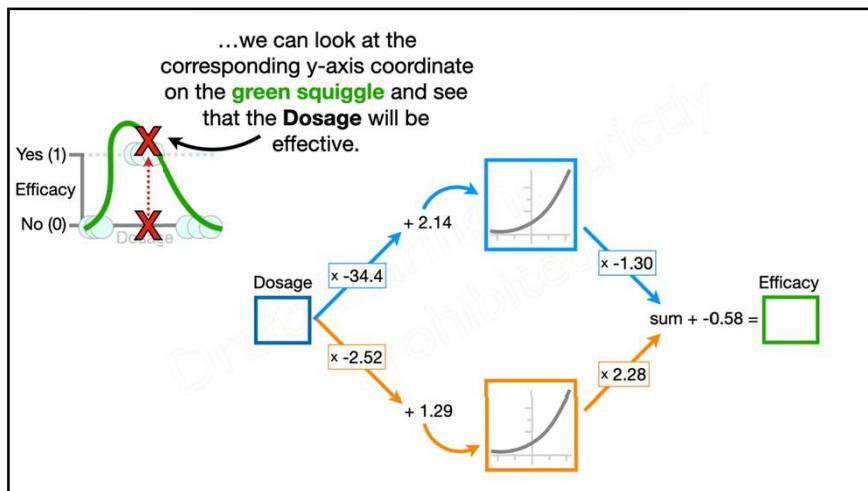
So we see that fitting a **Neural Network** to data gives us different parameter estimates on the connections.

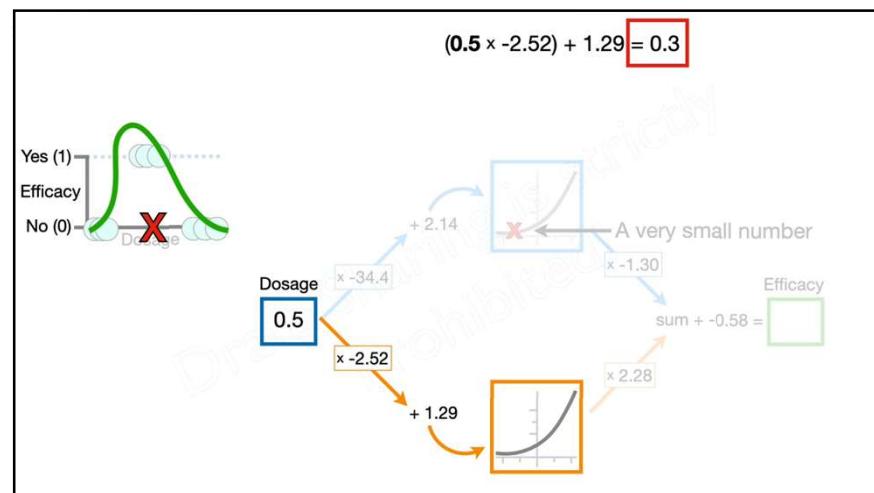
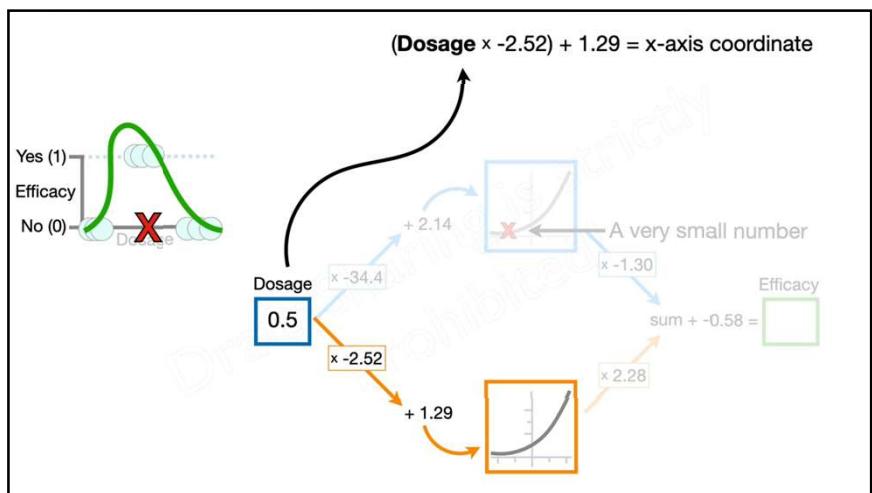
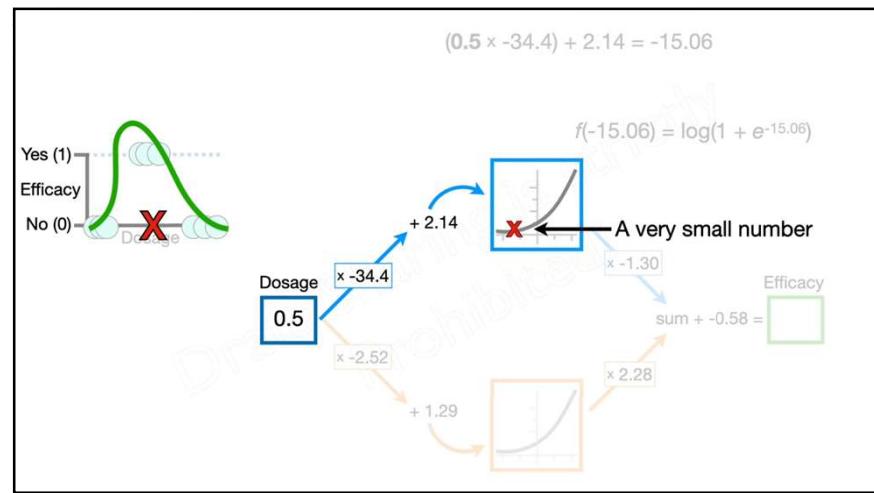
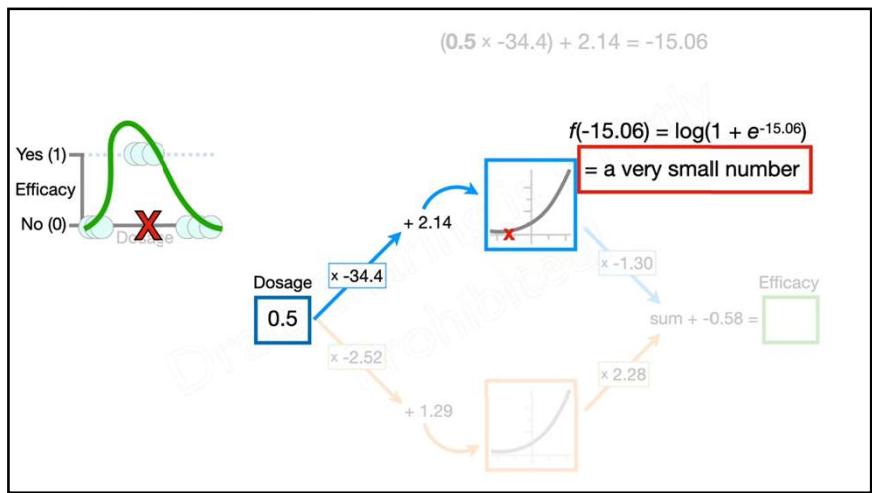


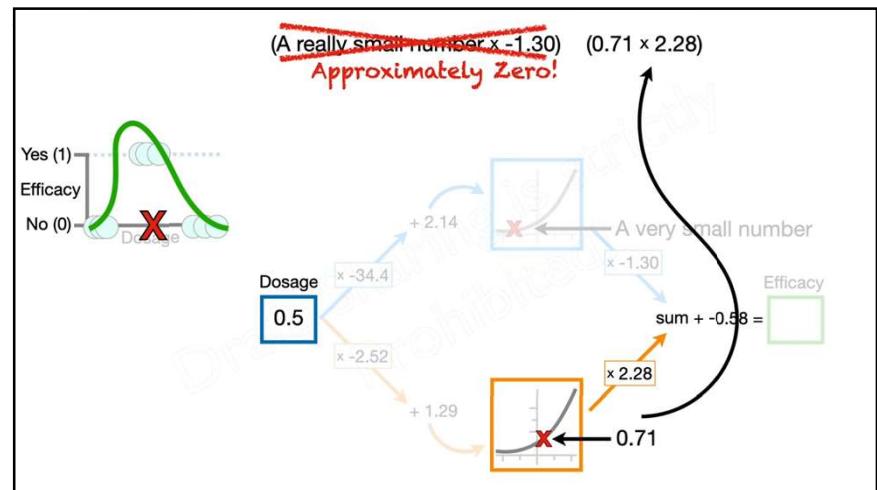
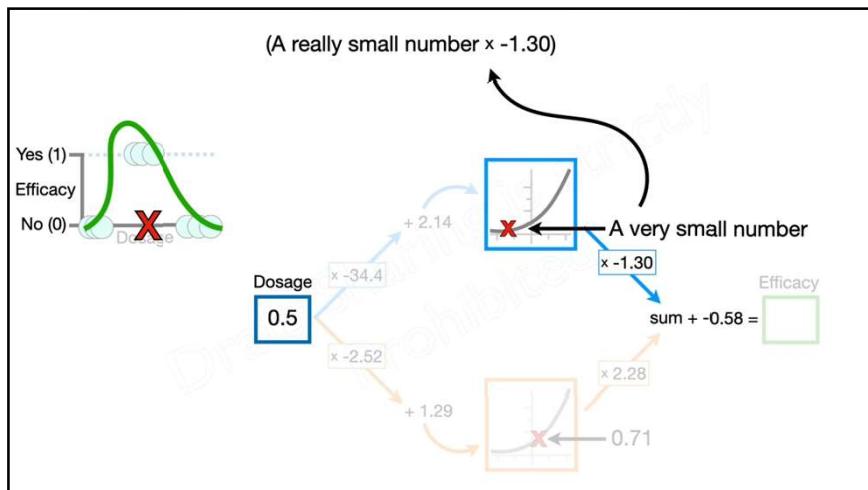
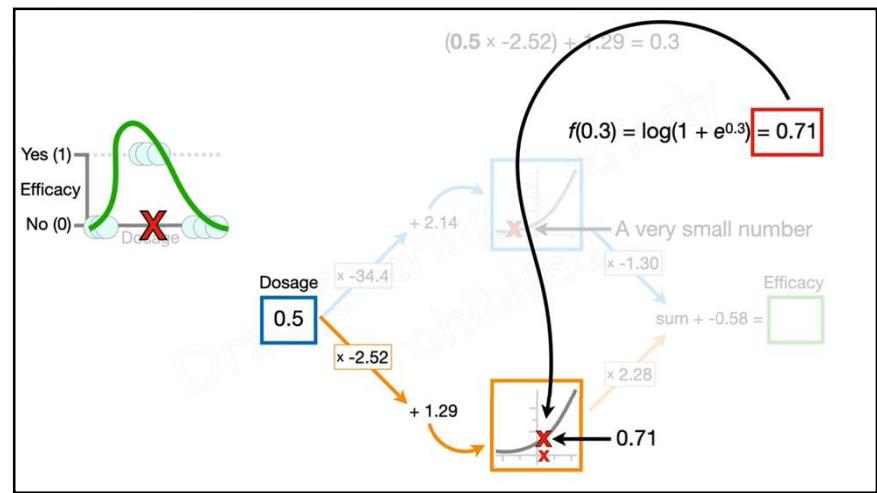
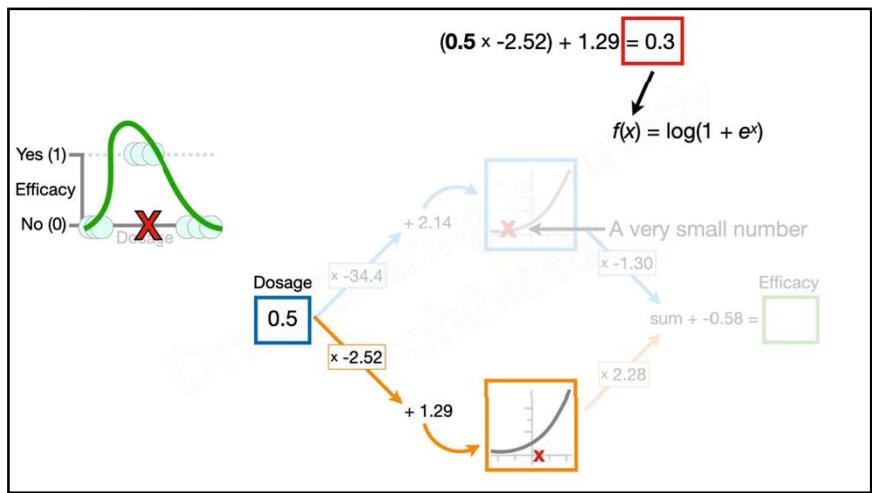


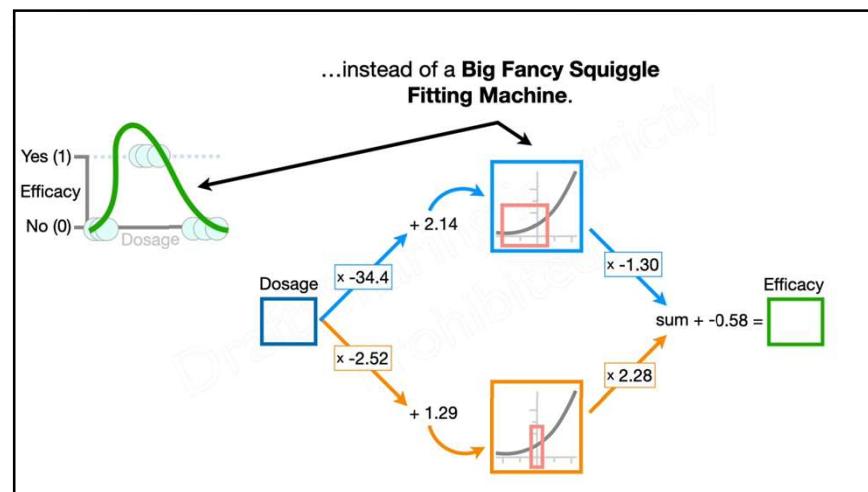
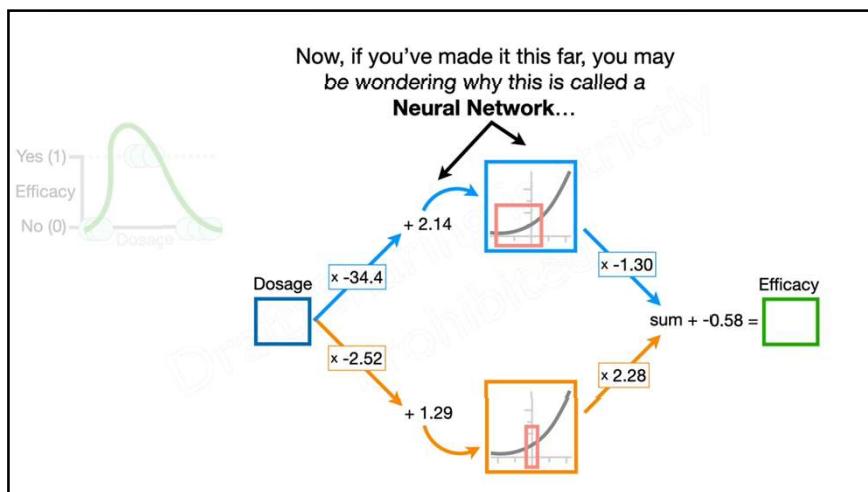
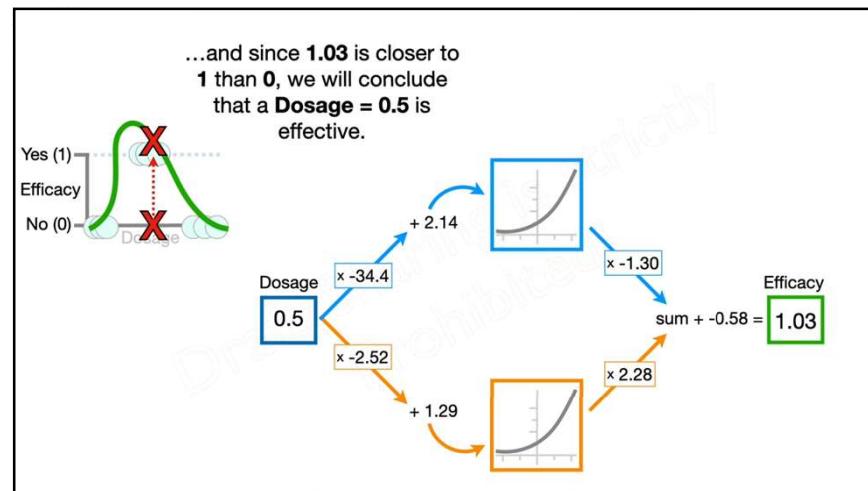
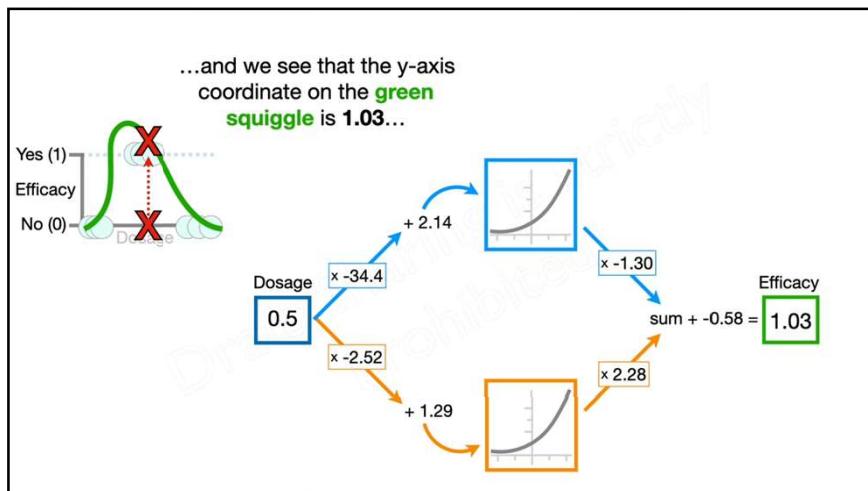




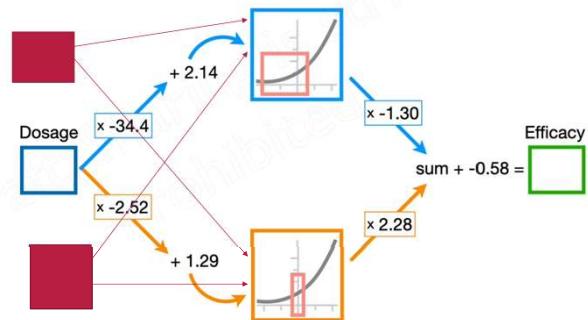




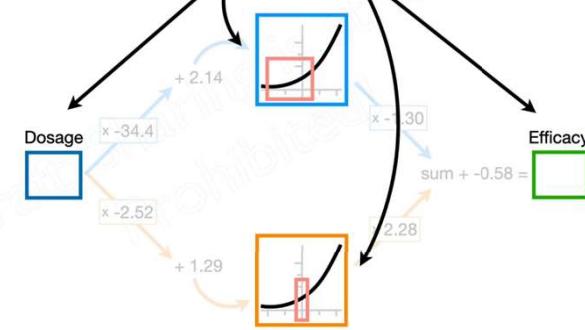




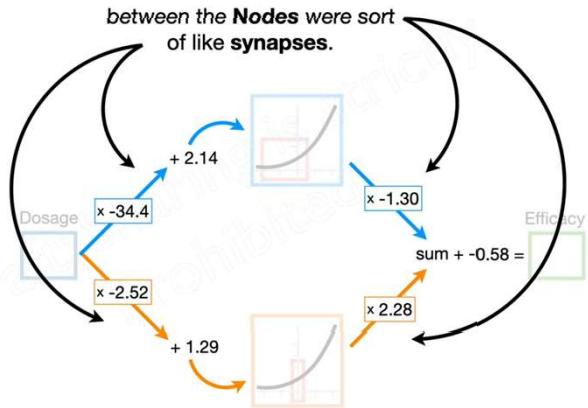
The reason is that way back in the **1940s and 50s** when **Neural Networks** were invented...



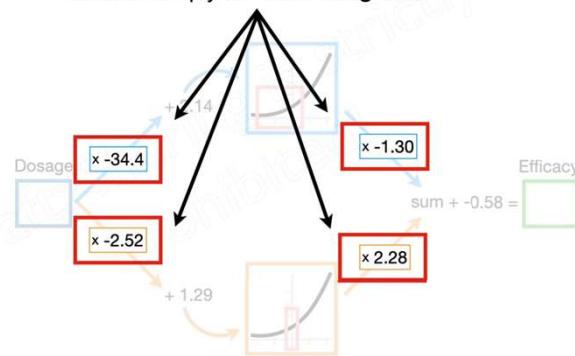
...they thought the **Nodes** were vaguely like **Neurons**...

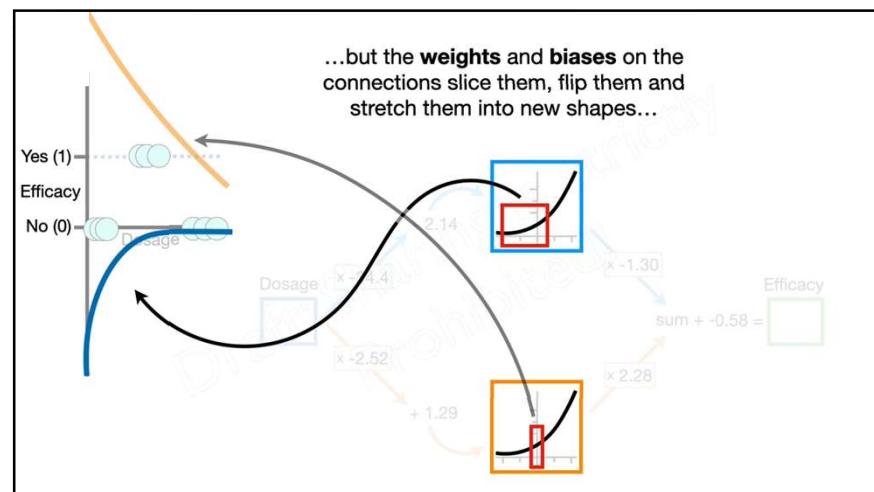
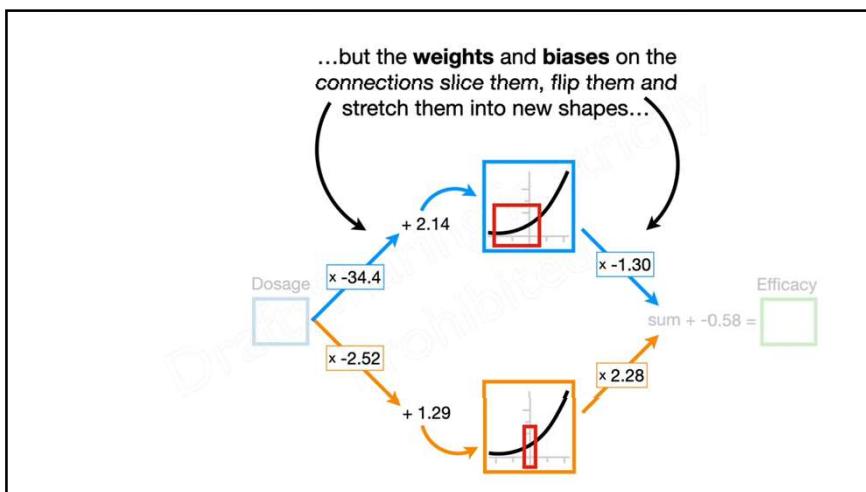
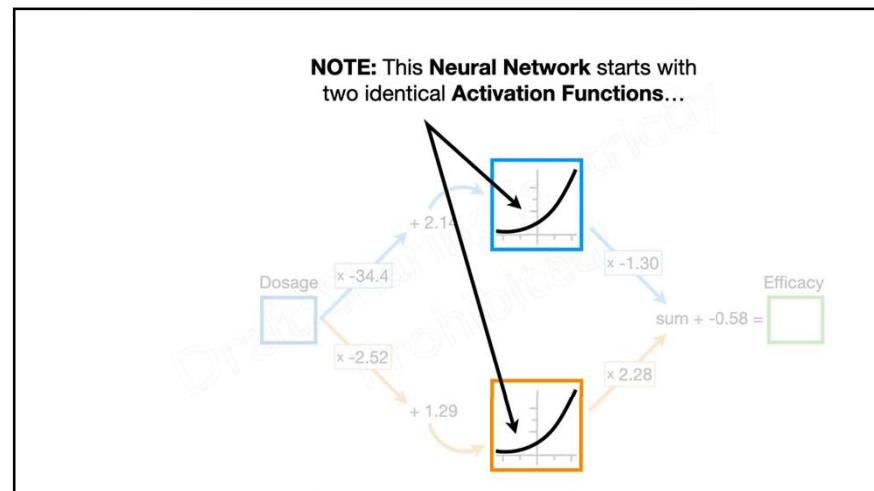
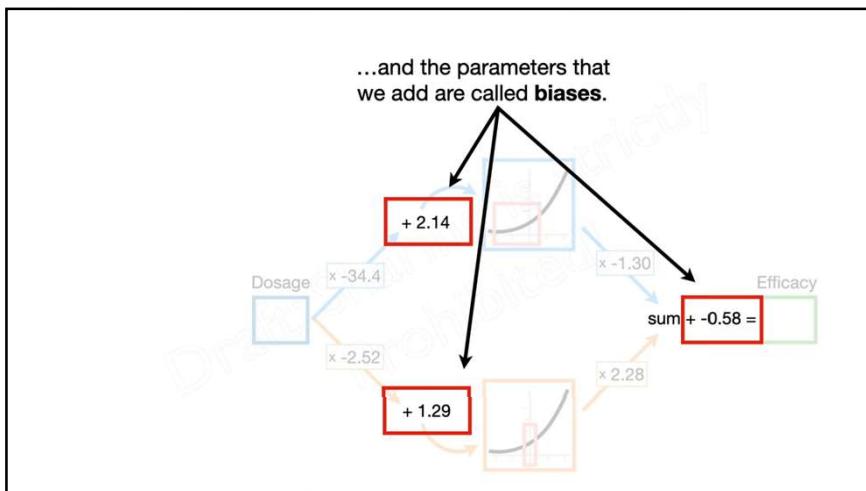


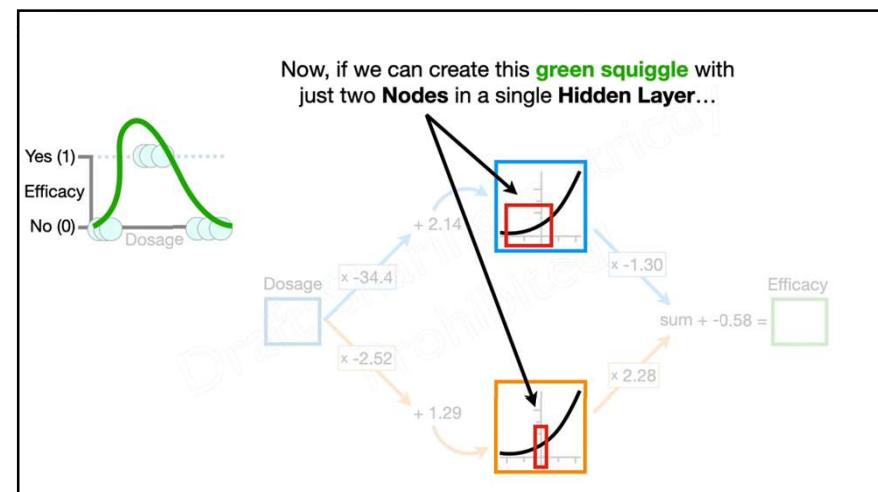
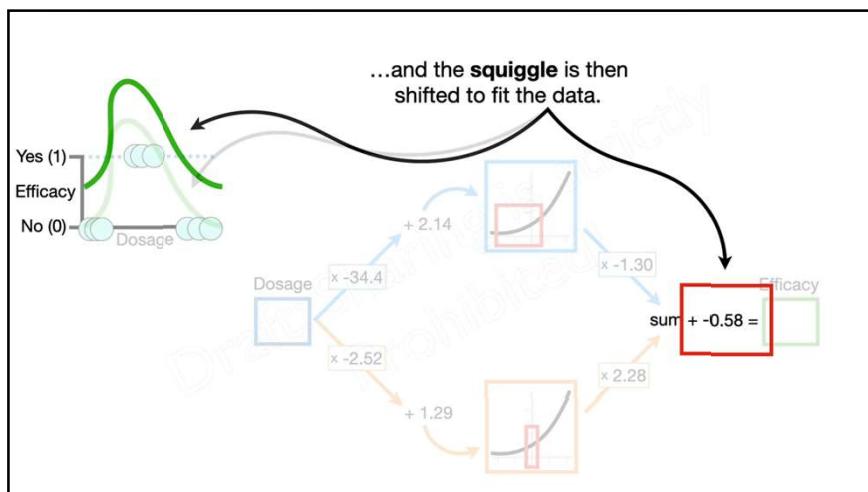
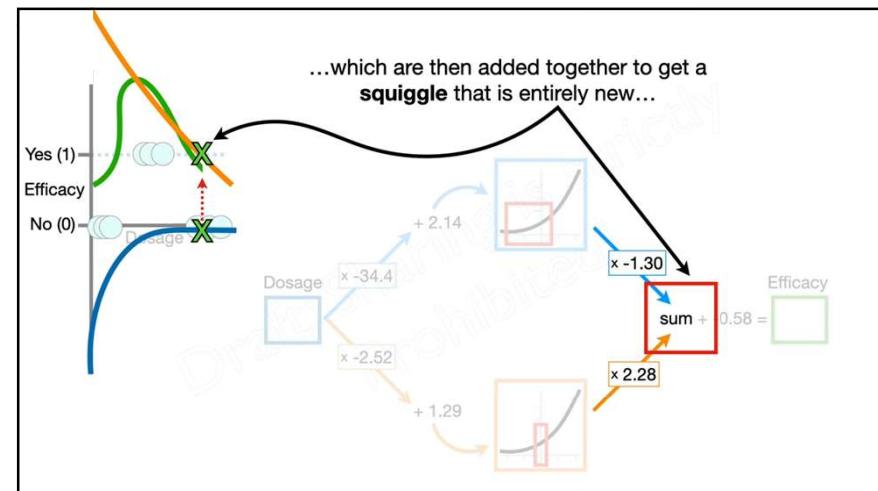
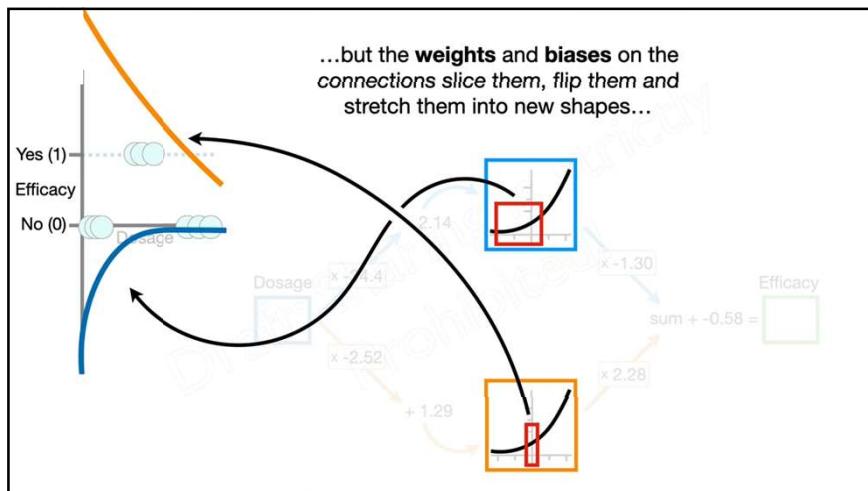
...and the **connections** between the **Nodes** were sort of like **synapses**.



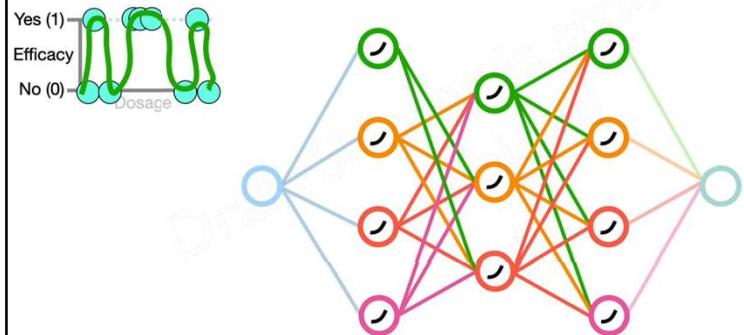
NOTE: Whether or not you call it a **Squiggle Fitting Machine**, the parameters that we multiply are called **weights**...



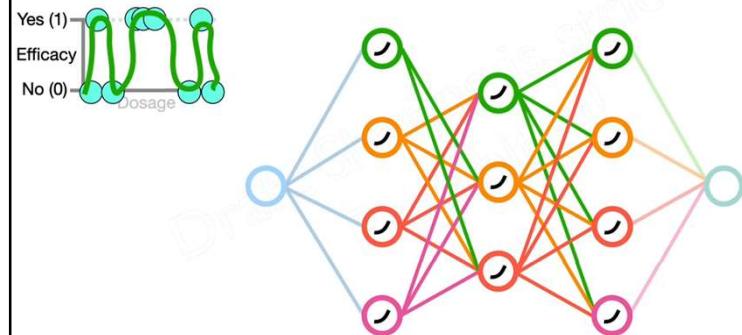




...just imagine what types of **green squiggles**
we could fit with more **Hidden Layers** and
more **Nodes** in each **Hidden Layer**.



In theory, **Neural Networks** can fit a **green squiggle** to just about any dataset, no matter how complicated!



THANK YOU!